



CSE2431 – Lecture Topic 2

Process (part 1)





Process (Part 1)

Instructor: Luan Duong, Ph.D.

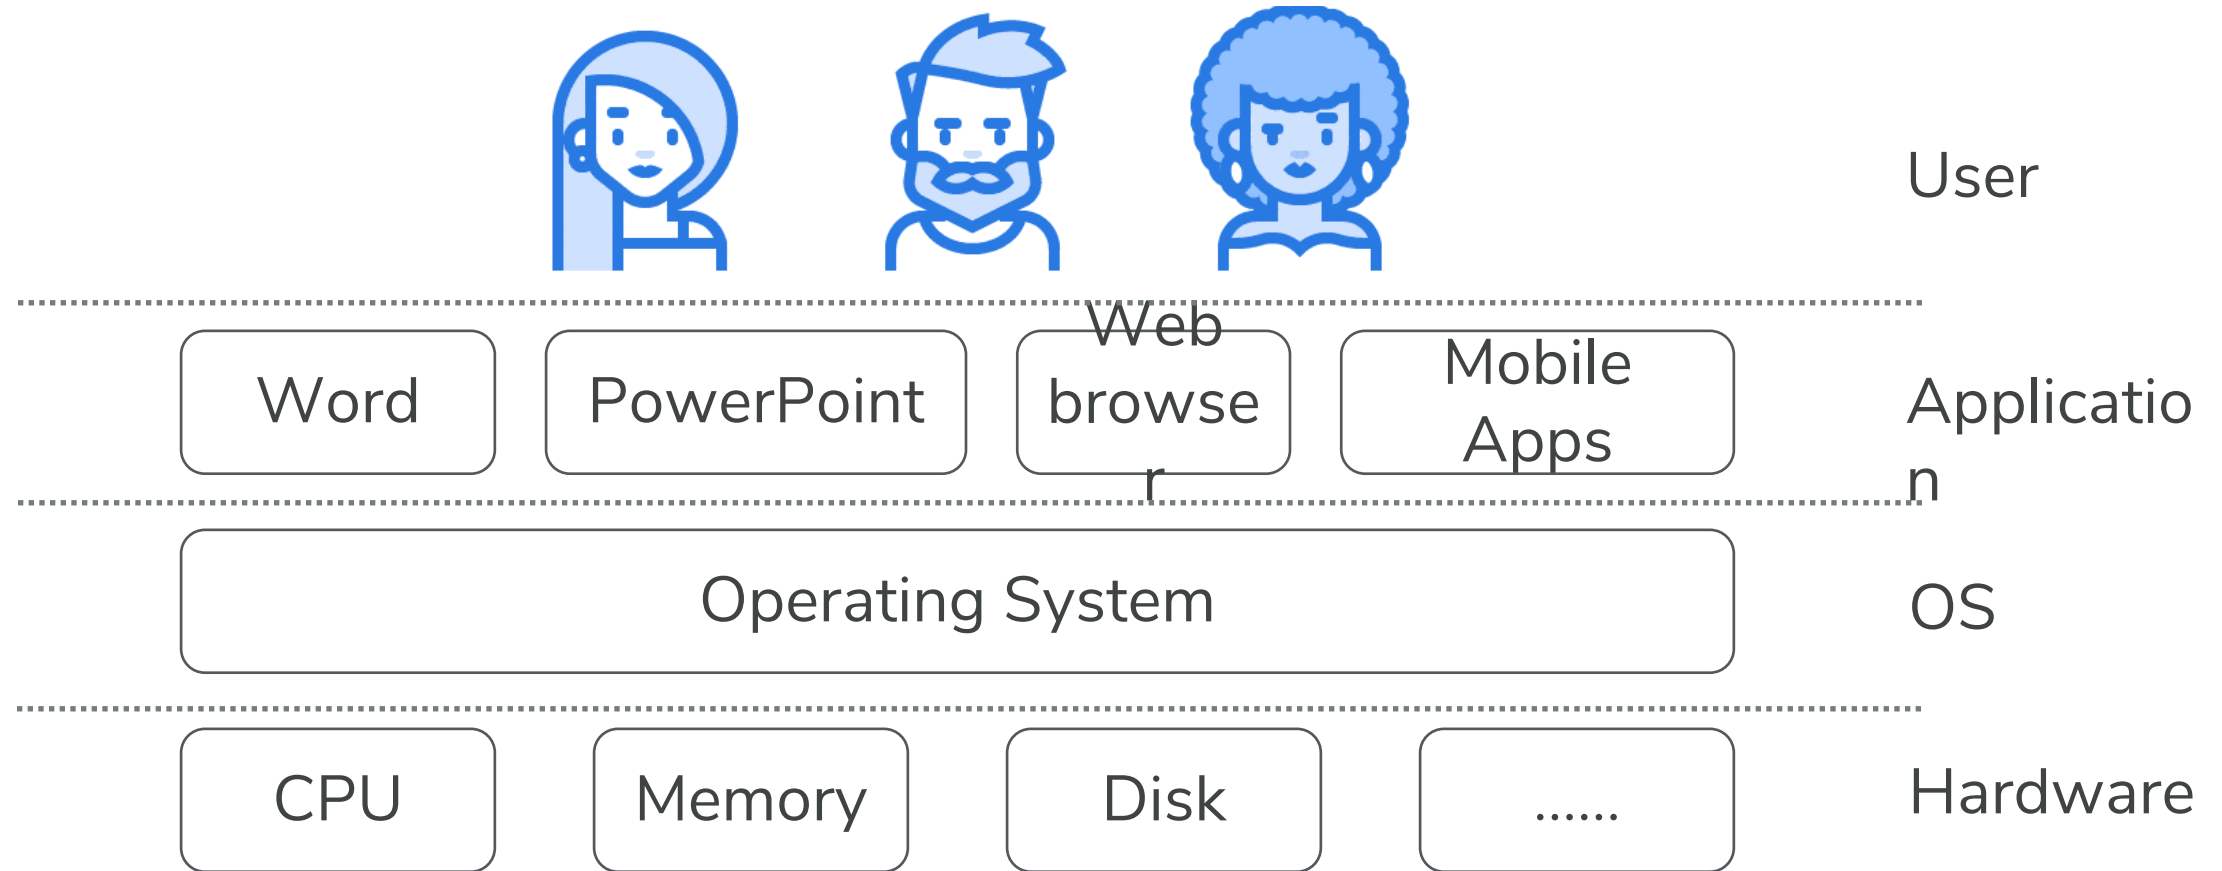
CSE 2431: Introduction to Operating Systems

Reading: Chapter 4-5 in required textbook

Lecture slides and materials adopted and referred from previously taught course by Dr. Yang Wang and Dr. Adam C. Champion

Last lecture

- Computer Architecture:



Last lecture

- Computer Architecture:

Operating System

Provide services

- Abstraction
- Convenience
- Standardization

Manage resources

- Allocation (CPUs, Memory, I/O devices)
- Reclamation (Voluntary at runtime, preemptive...)
- Protection (Protect from unauthorized access)
- Virtualization (Virtual memory, timeshared CPU)

Outline: Process

- What is a process?
- Process States; Process Control Block (PCB)
- Process Creation; *fork* command
- Process Memory Layout
- Process Scheduling
- Context Switch
- Process Operations
- Inter-Process Communication
- Client-Server Communication

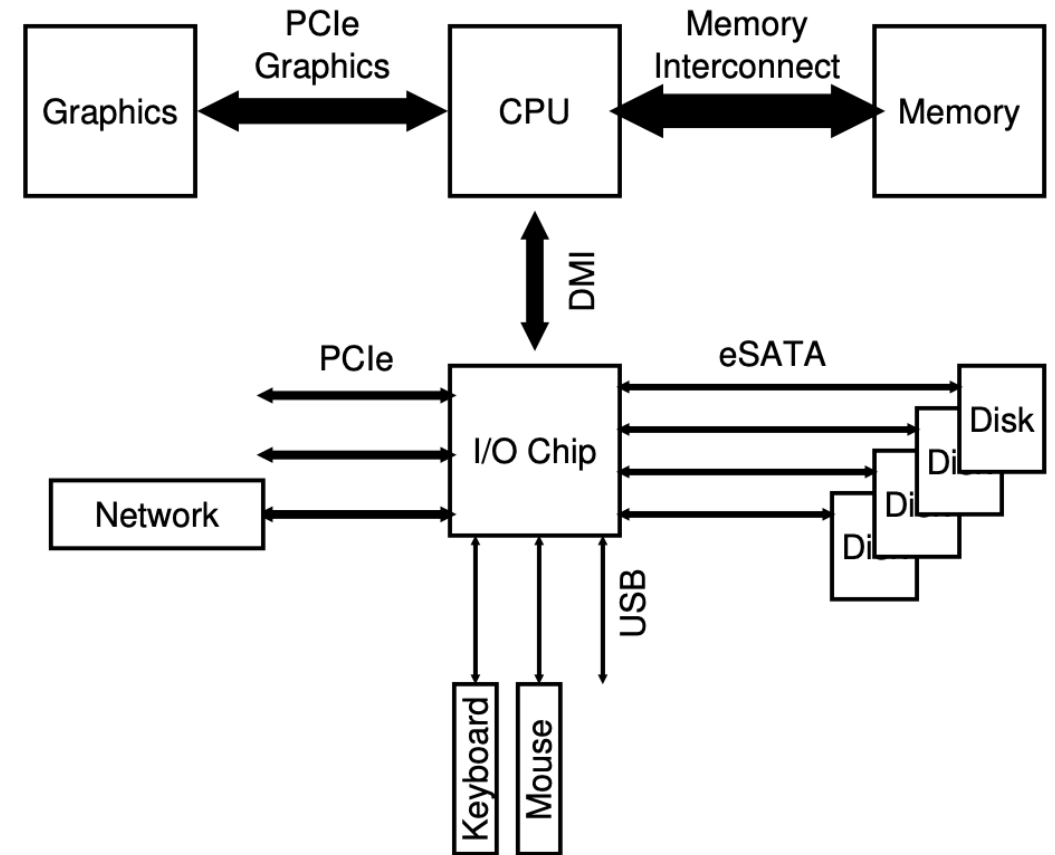
What is a process?

- Back to a question from previous lecture:
 - We only have one CPU, how could we run multiple programs or commands at the same time?
 - How to provide the illusion of many CPUs?

Virtualization

- So how can we call/abstract a thing provided by the OS through virtualization?

A Process



What is a process?

- Users launch programs
 - Many users can launch the same programs
 - One user can launch many instances of the same programs
- Processes: an executing 'piece' of program
- Do you think of any real-life analogies? → Class discussion

What is a process?

- A **process** is simply a **running program**; at **any instant in time**. It is also separated from other instances.
 - On batch system, refer to *jobs*
 - On interactive system (today OS), refer to *processes*
- Process can launch other processes or can be launched by others.
- Process \neq Program:
 - A program is **static**, while a process is **dynamic**.
- A process includes:
 - A program counter
 - Stack
 - Data section

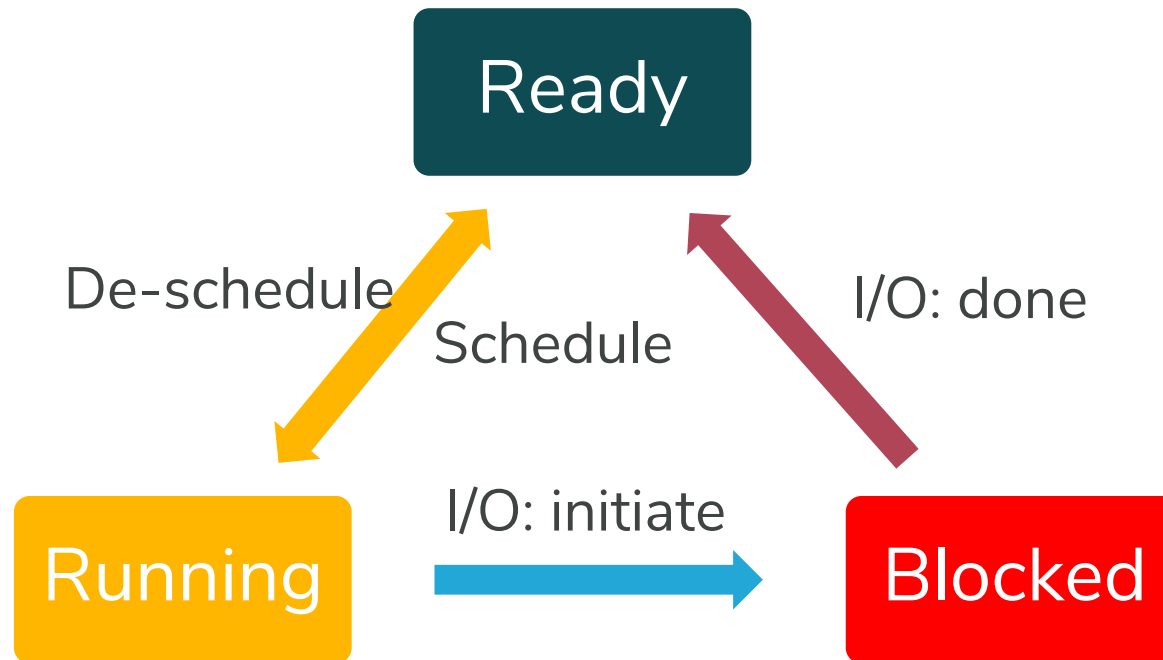
How can OS perform virtualization?

- Time sharing (and Space sharing)
 - “Allowing the resource to be used for a little while by one entity, and then a little while by another, and so forth” [Operating Systems: Three Easy Pieces]
 - Each process becomes slower, but users usually cannot tell, because computers are much faster than human beings!
 - Counterpart of time sharing is space sharing, where a resource (CPU) is divided (in space/memory) among those who wish to use it.
- Process operation
 - Create
 - Destroy
 - Wait
 - Miscellaneous Control
 - Status (i.e. Process state)



Process status

- Running: a process is being executed by a CPU
- Ready: a process is ready to run but is not running
- Blocked: a process is waiting on some event to take place



Process status: Examples

- Running: a process is being executed by a CPU
- Ready: a process is ready to run but is not running
- Blocked: a process is waiting on some event to take place

Time	Process[0]	Process[1]	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process[0] now done
4	[Destroyed]	Running	
5		Blocked	I/O initiate
...			(keyboard interrupt → run other processes)
11		Ready	I/O done
12		Running	
13		[Destroyed]	Process[1] now done

Process status: Examples

- Running: a process is being executed by a CPU
- Ready: a process is ready to run but is not running
- Blocked: a process is waiting on some event to take place

Time	Process[0]	Process[1]	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process[0] initiates I/O
4	Blocked	Running	Process[0] is blocked, so Process[1] runs
5	Blocked	Running	
6	Ready	Running	I/O done
7	Ready	Running	Process[1] now done
8	Running	[Destroyed]	Process[0] now done

Process status: PCB

- How to maintain a lot of processes' status? → Process Control Block (PCB)
 - Process state
 - Process identification (PID)
 - Program counter
 - CPU registers
 - CPU scheduling info
 - Memory-management info
 - Accounting info
 - I/O status info
 - PID of parent process

Process status: PCB

- How to maintain a lot of processes' status? → Process Control Block (PCB)

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;
    int esp;
    int ebx;
    int ecx;
    int edx;
    int esi;
    int edi;
    int ebp;
};
```

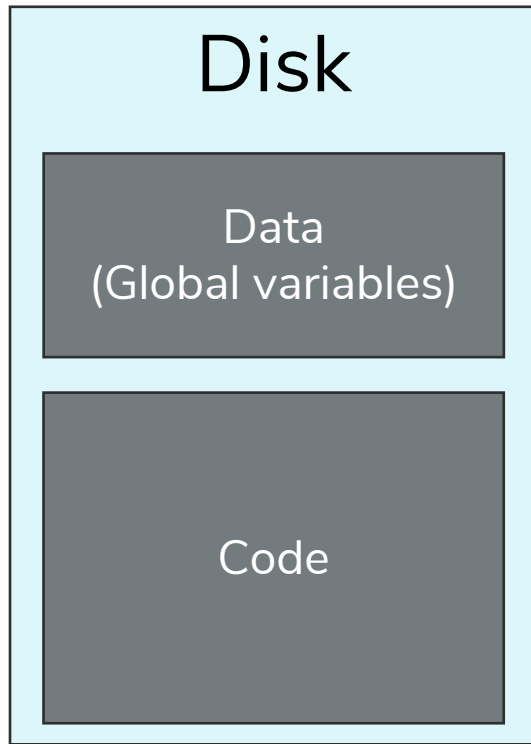
```
// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                // Start of process memory
    uint sz;                  // Size of process memory
    char *kstack;             // Bottom of kernel stack
                                // for this process
    enum proc_state state;    // Process state
    int pid;                  // Process ID
    struct proc *parent;      // Parent process
    void *chan;               // If !zero, sleeping on chan
    int killed;               // If !zero, has been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;         // Current directory
    struct context context;    // Switch here to run process
    struct trapframe *tf;     // Trap frame for the
                                // current interrupt
};
```

Example of a PCB: The xv6 Process Structure [Operating Systems: Three Easy Pieces]

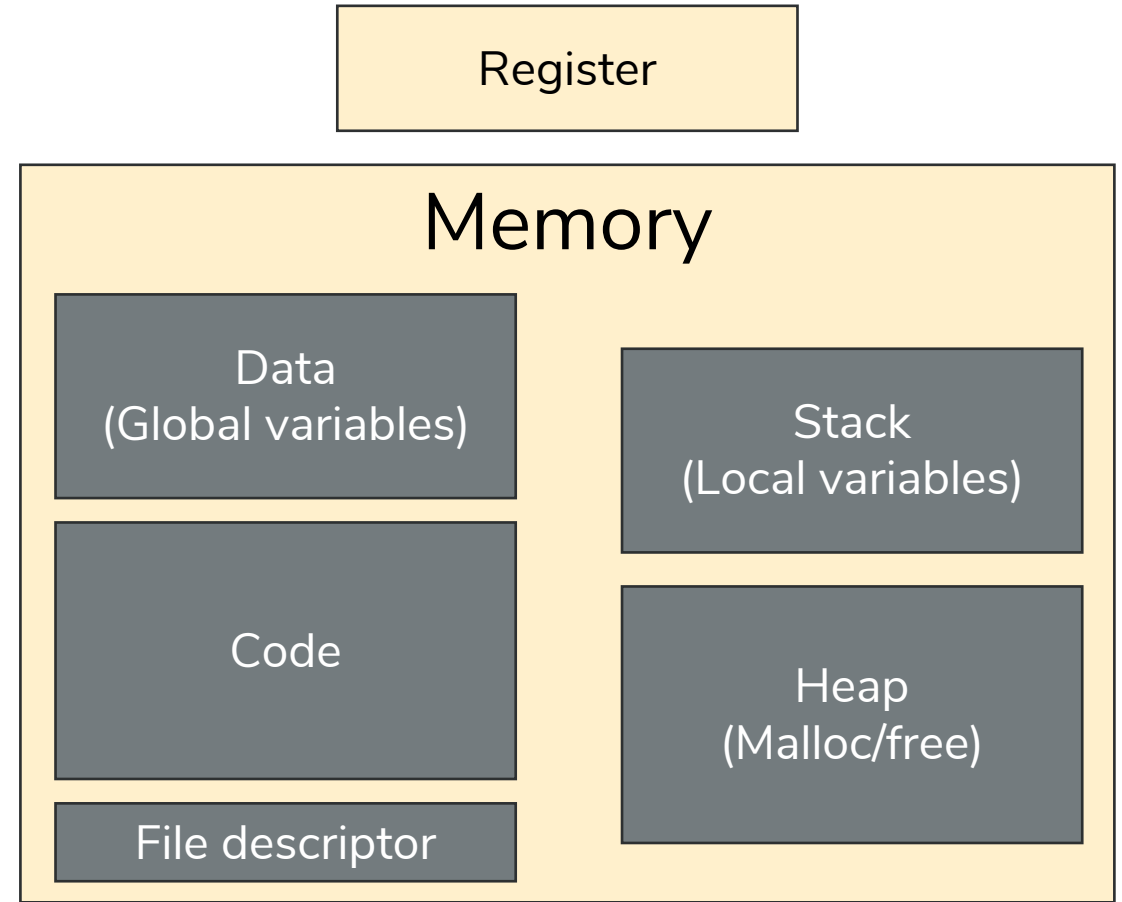
Process Creation

- Process Creation:



After your program is compiled, it usually consists of two parts

When a process is created, these parts will be loaded into memory (maybe partially)

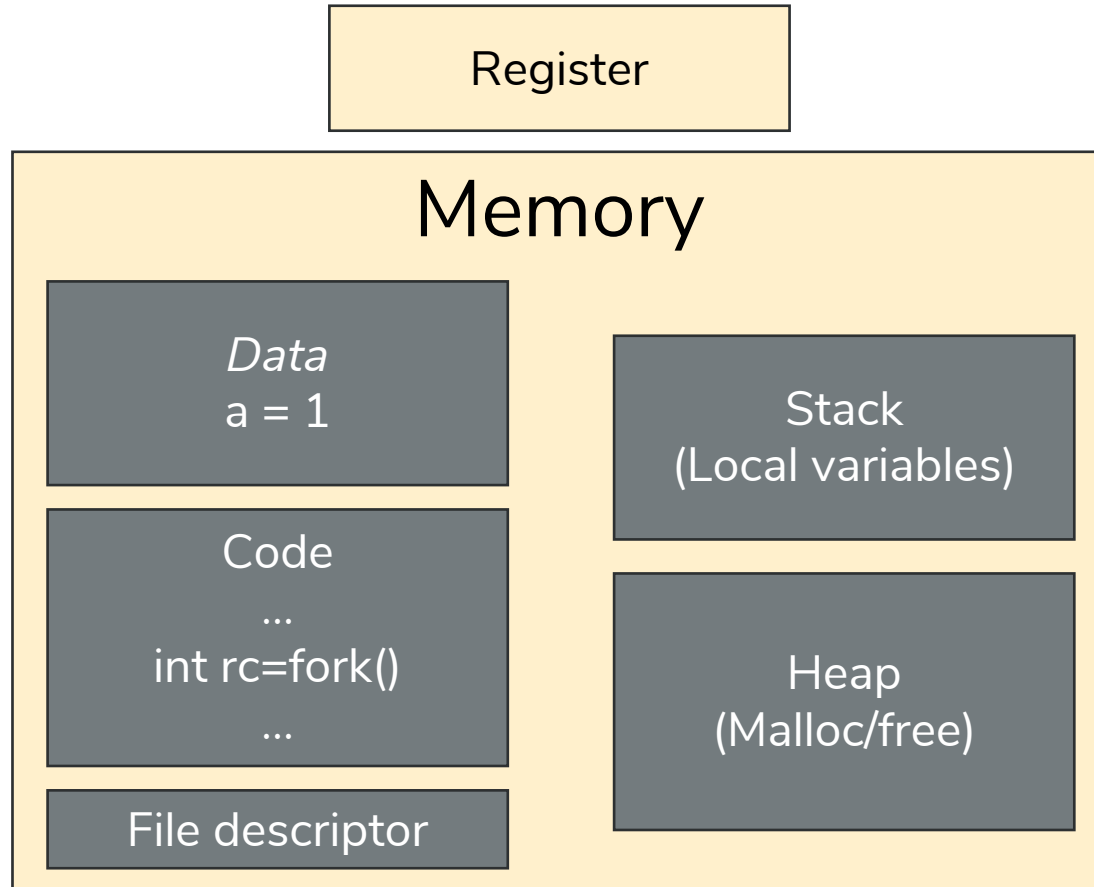


OS creates additional data structures for a process

Process Creation through System calls

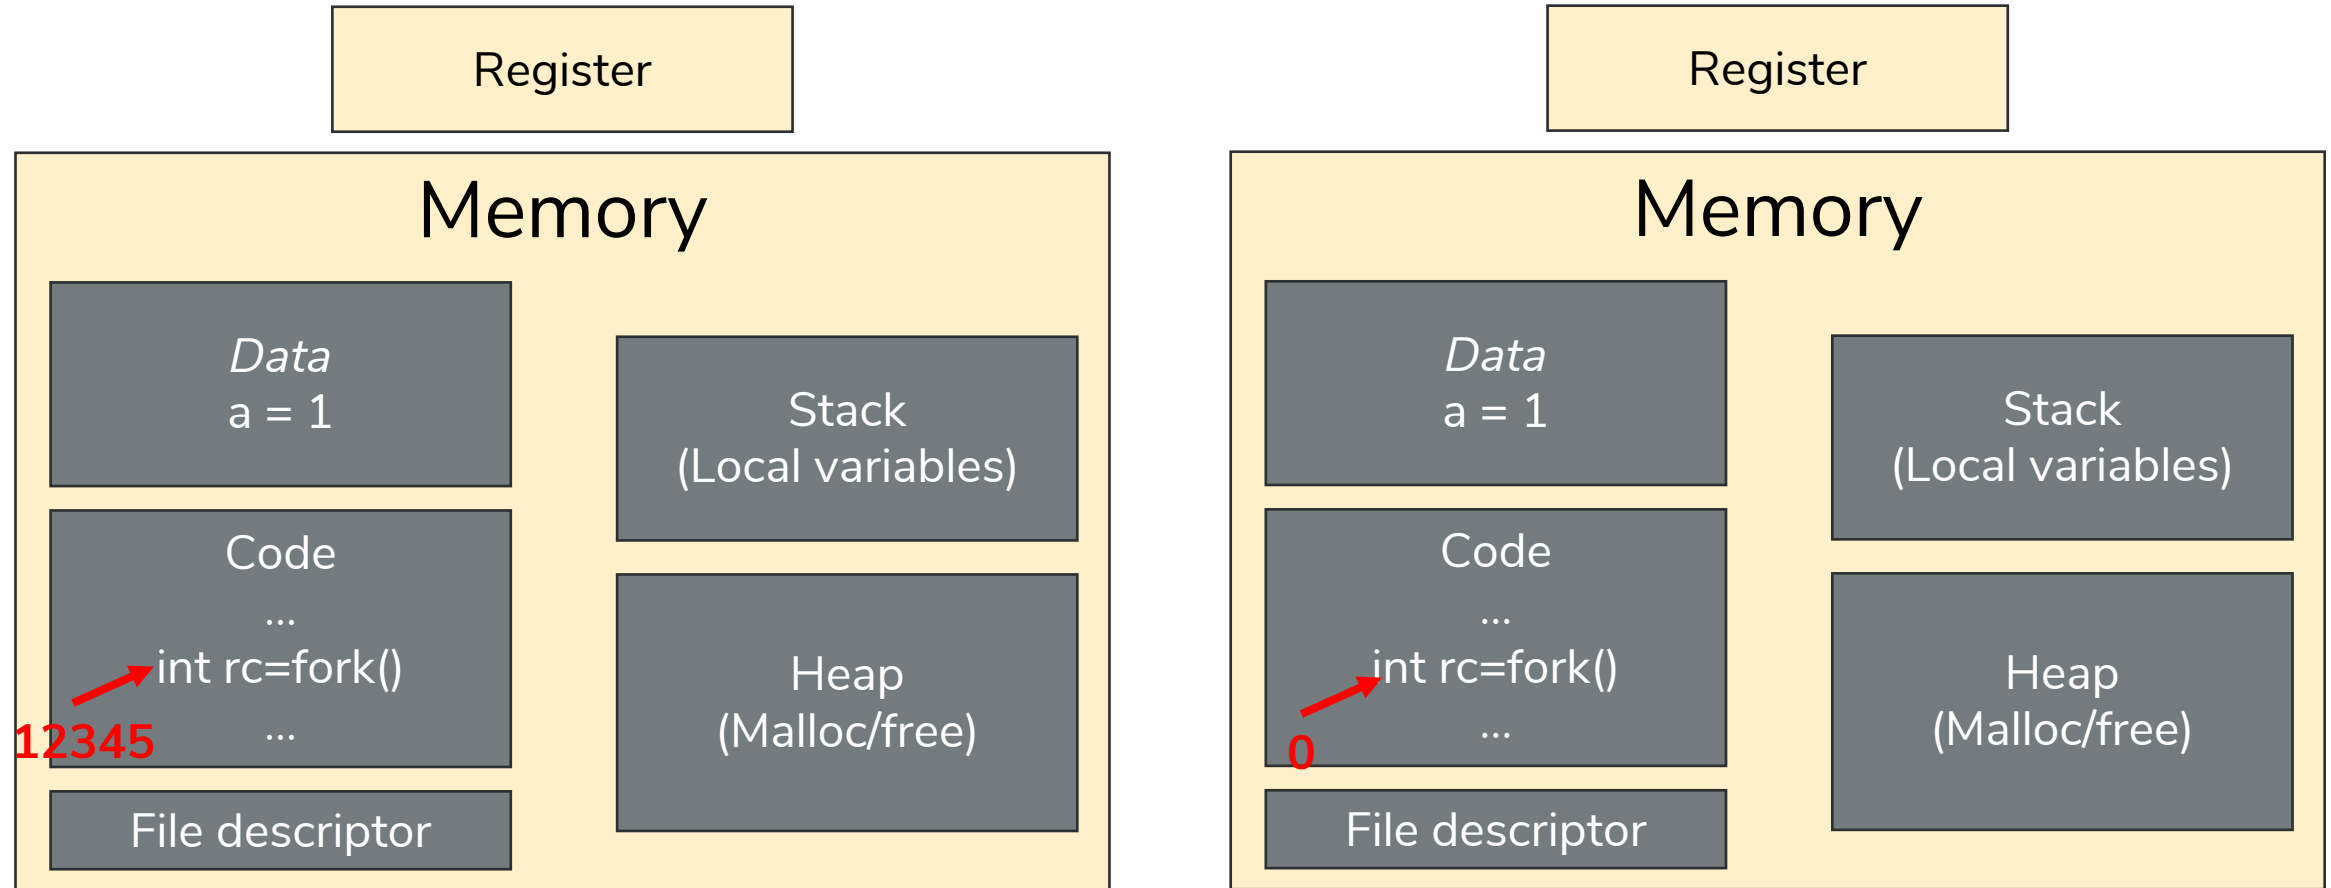
- ***fork***: duplicate the process that is calling fork
- ***exec***: run a new program (the new process will replace the current process)
- *fork* and *exec* are often used in combination to create a new process
- Other useful system calls: wait, pid, ...

Process Creation: *fork()*



Parent process --- the process that is calling fork

Process Creation: *fork()*



Parent process --- the process that is calling fork

Child process

Parent and child processes are almost identical except one thing: return value of fork. For parent, fork returns the process ID of the child. For child, fork returns 0. Why?

Process Creation: *fork()*

- First: A parent process sometimes needs to control its child processes
 - For example: wait for a child process to terminate; kill a child process; ...
 - So a parent process needs the process ID of its child
- Second: we need a way to differentiate parent and child processes
 - This means the return value to the child should not be a valid process ID: that is why we give it 0 (negative return value usually indicates error in C)
- Solution: return 0 to child; return child process id to parent

Process Creation: *fork()* example

```
int rc = fork();
if (rc < 0) {
    // fork failed
    fprintf(stderr, "fork failed\n");
    exit(1);
} else if (rc == 0) {
    // child (new process)
    printf("child (pid:%d)\n", (int) getpid());
} else {
    // parent goes down this path (main)
    printf("parent of %d (pid:%d)\n",
           rc, (int) getpid());
}
```

Process Creation: *fork()* – Important facts

- Parent and child processes **do NOT share memory**
- Parent's execution after *fork* **will NOT** affect child, and vice versa (which means parent and child can either execute concurrently or parent waits until child terminates it depends)
- Parent and child are executed in parallel and can be executed **in any order**.
- In Unix system:
 - All resources shared (i.e. child is a clone)
 - `Execve()` system call used to replace process' memory with a new program.

Process Creation: *fork()* – Some exercise

- What is the output of this code?
- Naturally, you will think it will output this:

```
prompt> ./p1
hello (pid:29146)
child (pid:29147)
parent of 29147 (pid:29146)
prompt>
```

- But another case can happen. Why?

```
int main(int argc, char *argv[]) {
    printf("hello (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("child (pid:%d)\n", (int) getpid());
    } else {
        // parent goes down this path (main)
        printf("parent of %d (pid:%d)\n",
               rc, (int) getpid());
    }
    return 0;
}
```

Let's do an exercise


```
int num = 3;

int main(int argc, char**argv) {
    printf("start num is %d\n", num); num++;
    int rc = fork();
    if(rc>0)
    {
        num+=2;
        printf("Parent's num is %d\n", num);
    }
    else if(rc==0)
    {
        num+=3;
        printf("Child's num is %d\n", num);
    }
    num+=4;
    printf("end num is %d\n", num);
}
```

What is the output of this program?

Let's do an exercise

```
int num = 3;
```

```
int main(int argc, char**argv) {  
    printf("start num is %d\n", num);   
    num++;  
    int rc = fork();  
    if(rc>0)  
    {  
        num+=2;  
        printf("Parent's num is %d\n", num);  
    }  
    else if(rc==0)  
    {  
        num+=3;  
        printf("Child's num is %d\n", num);  
    }  
    num+=4;  
    printf("end num is %d\n", num);  
}
```

One process (num = 3):

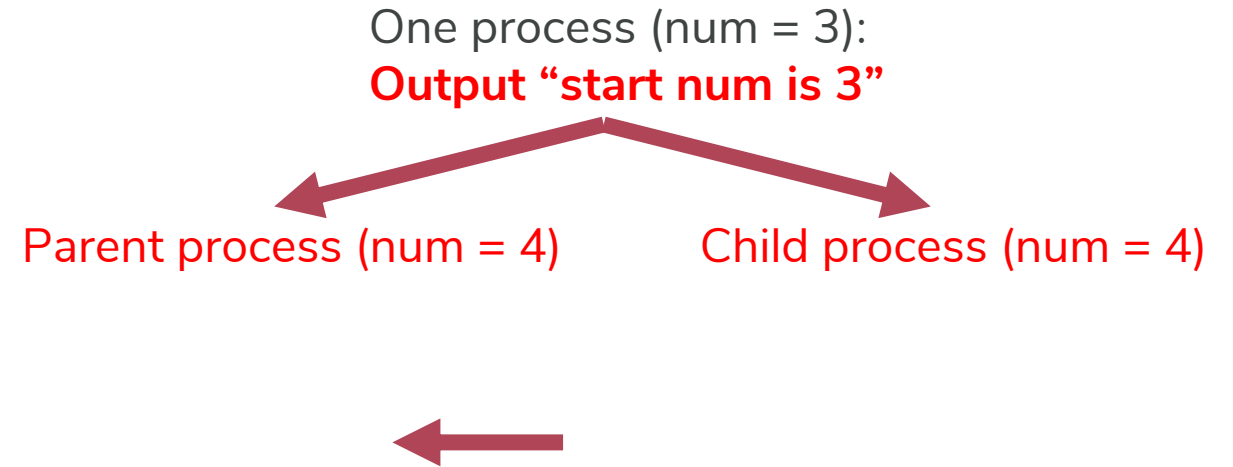
Output “start num is 3”



Let's do an exercise


```
int num = 3;

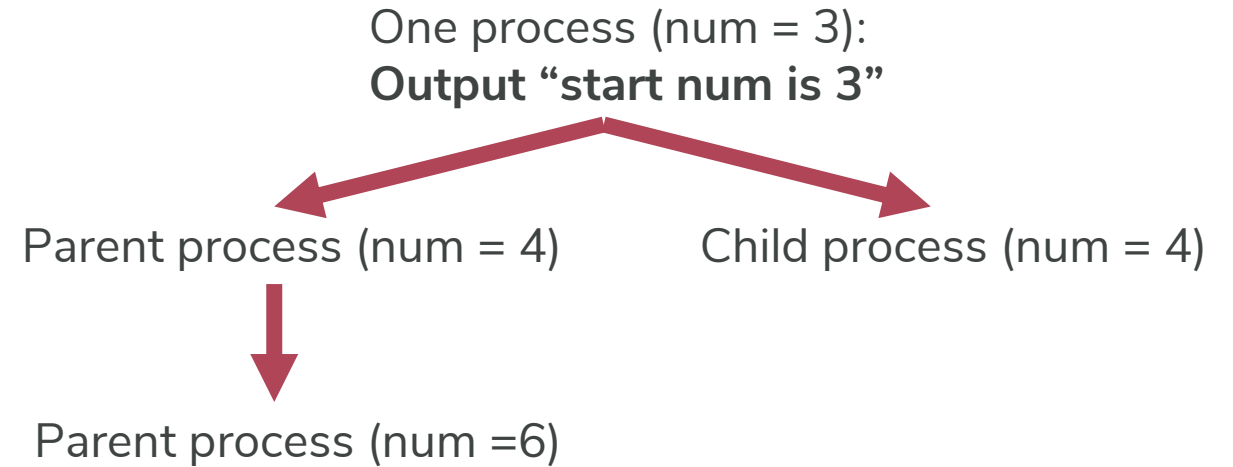
int main(int argc, char**argv) {
    printf("start num is %d\n", num);
    num++;
    int rc = fork();
    if(rc>0)
    {
        num+=2;
        printf("Parent's num is %d\n", num);
    }
    else if(rc==0)
    {
        num+=3;
        printf("Child's num is %d\n", num);
    }
    num+=4;
    printf("end num is %d\n", num);
}
```



Let's do an exercise

```
int num = 3;
```

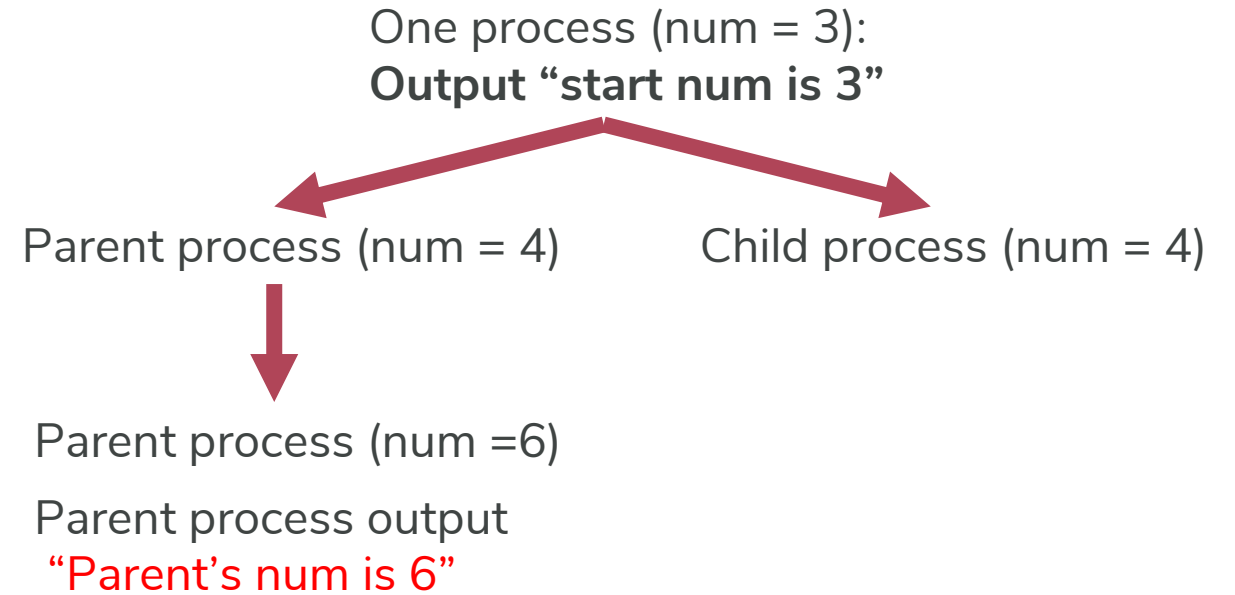
```
int main(int argc, char**argv) {  
    printf("start num is %d\n", num);  
    num++;  
    int rc = fork();  
    if(rc>0)  
    {  
        num+=2;   
        printf("Parent's num is %d\n", num);  
    }  
    else if(rc==0)  
    {  
        num+=3;  
        printf("Child's num is %d\n", num);  
    }  
    num+=4;  
    printf("end num is %d\n", num);  
}
```



Let's do an exercise

```
int num = 3;
```

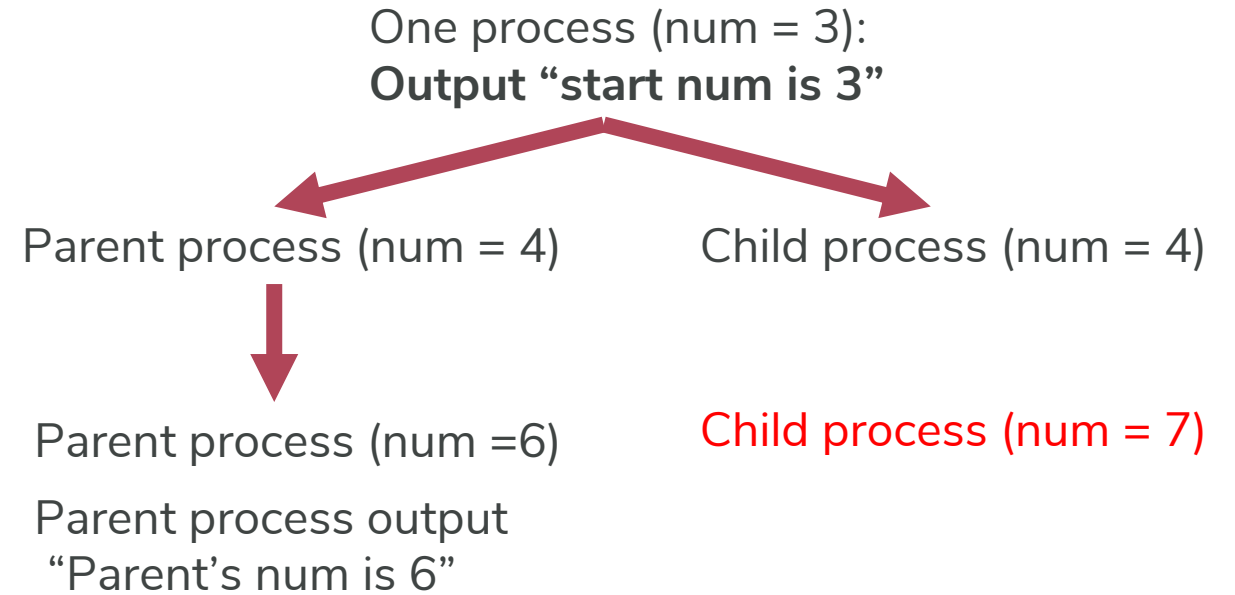
```
int main(int argc, char**argv) {  
    printf("start num is %d\n", num);  
    num++;  
    int rc = fork();  
    if(rc>0)  
    {  
        num+=2;  
        printf("Parent's num is %d\n", num);  
    }  
    else if(rc==0)  
    {  
        num+=3;  
        printf("Child's num is %d\n", num);  
    }  
    num+=4;  
    printf("end num is %d\n", num);  
}
```



Let's do an exercise

```
int num = 3;

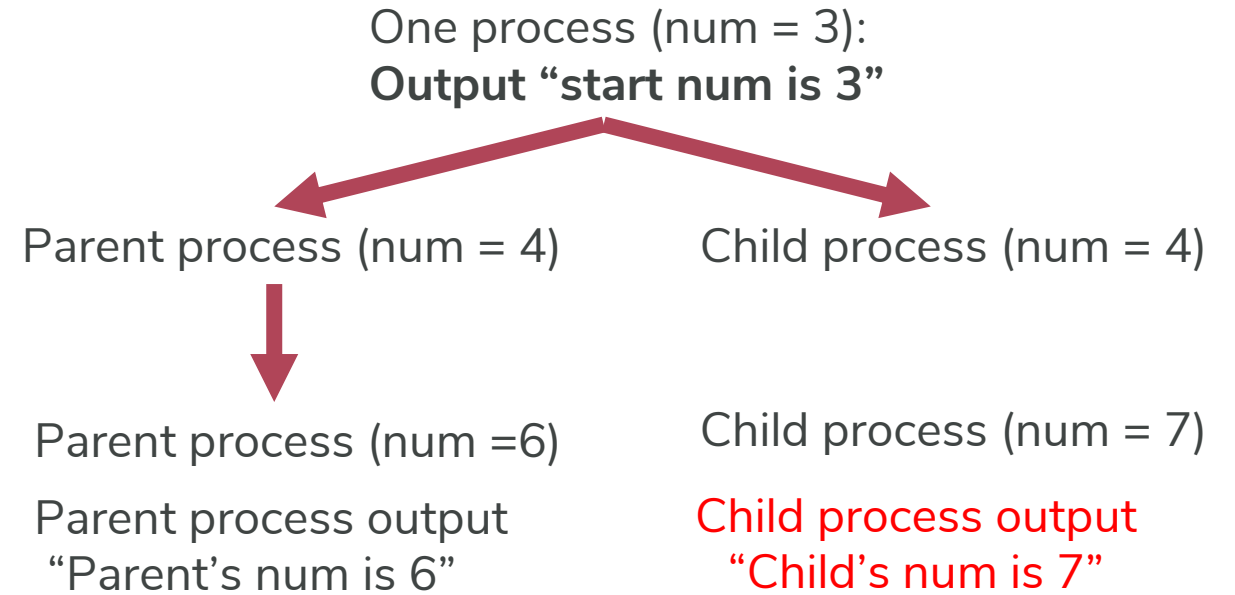
int main(int argc, char**argv) {
    printf("start num is %d\n", num);
    num++;
    int rc = fork();
    if(rc>0)
    {
        num+=2;
        printf("Parent's num is %d\n", num);
    }
    else if(rc==0)
    {
        num+=3;
        printf("Child's num is %d\n", num);
    }
    num+=4;
    printf("end num is %d\n", num);
}
```



Let's do an exercise

```
int num = 3;
```

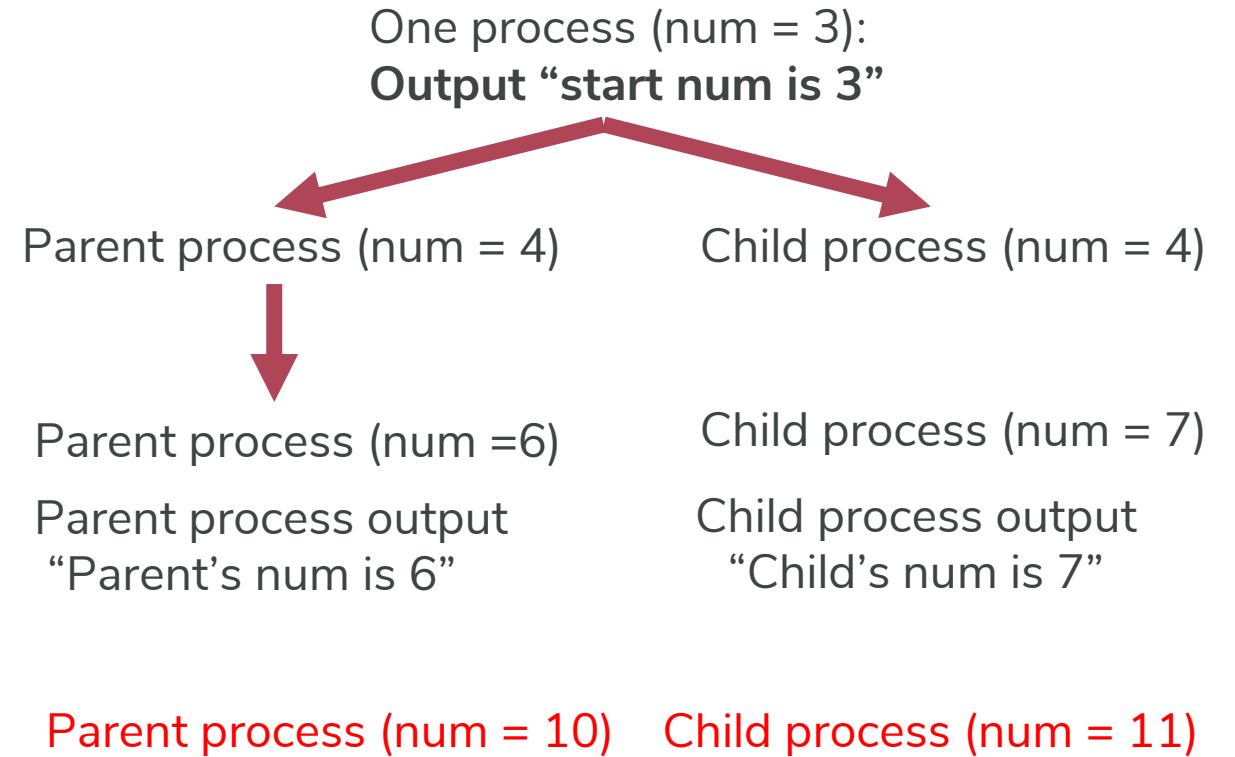
```
int main(int argc, char**argv) {  
    printf("start num is %d\n", num);  
    num++;  
    int rc = fork();  
    if(rc>0)  
    {  
        num+=2;  
        printf("Parent's num is %d\n", num);  
    }  
    else if(rc==0)  
    {  
        num+=3;  
        printf("Child's num is %d\n", num);  
    }  
    num+=4;  
    printf("end num is %d\n", num);  
}
```



Let's do an exercise

```
int num = 3;

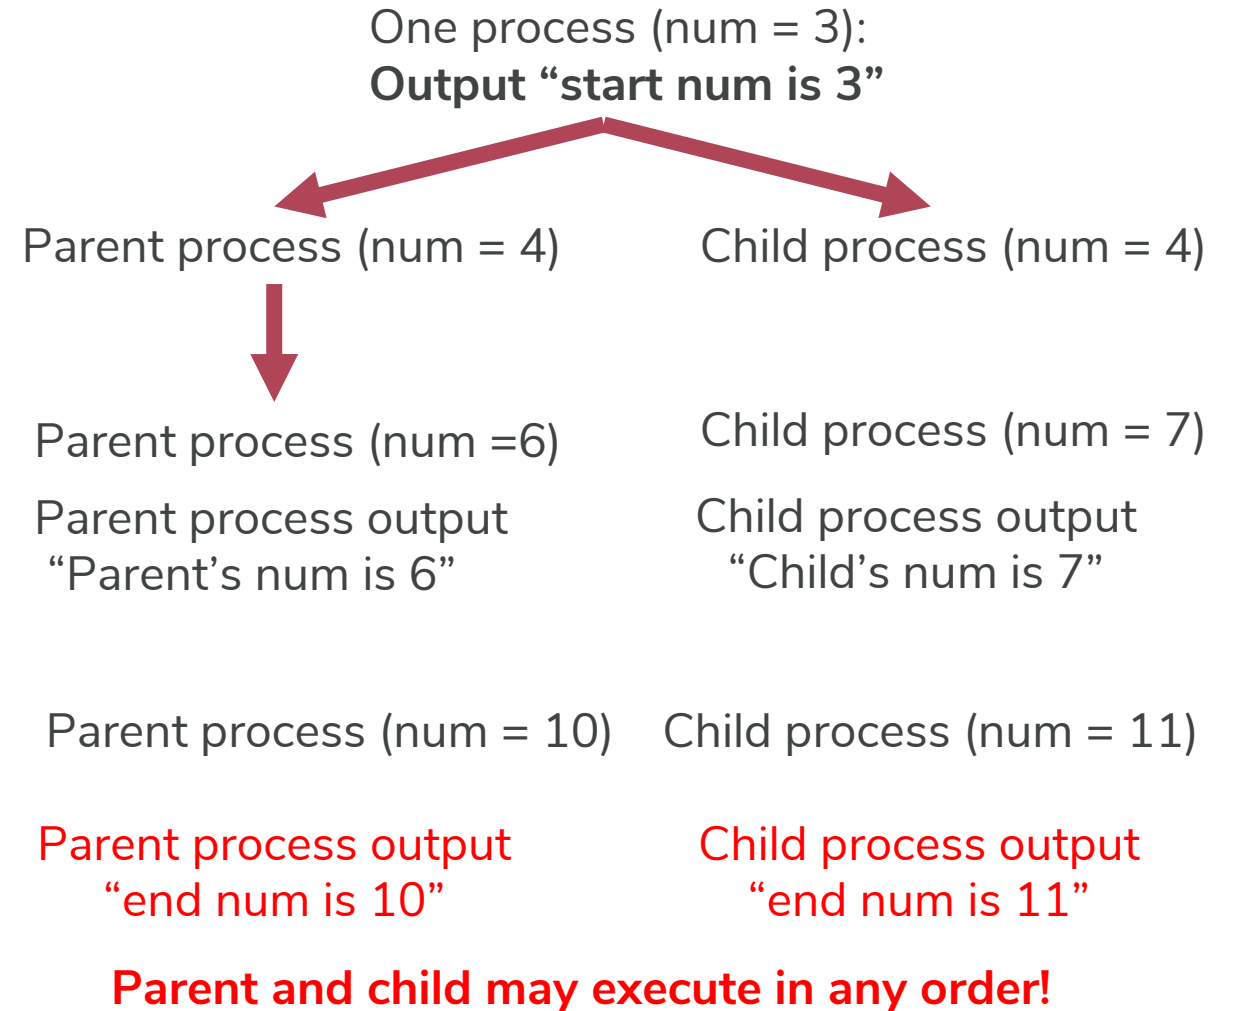
int main(int argc, char**argv) {
    printf("start num is %d\n", num);
    num++;
    int rc = fork();
    if(rc>0)
    {
        num+=2;
        printf("Parent's num is %d\n", num);
    }
    else if(rc==0)
    {
        num+=3;
        printf("Child's num is %d\n", num);
    }
    num+=4;
    printf("end num is %d\n", num);
}
```



Let's do an exercise

```
int num = 3;
```

```
int main(int argc, char**argv) {  
    printf("start num is %d\n", num);  
    num++;  
    int rc = fork();  
    if(rc>0)  
    {  
        num+=2;  
        printf("Parent's num is %d\n", num);  
    }  
    else if(rc==0)  
    {  
        num+=3;  
        printf("Child's num is %d\n", num);  
    }  
    num+=4;  
    printf("end num is %d\n", num);  
}
```



Let's do an exercise

Parent's execution (red)

```
int num = 3;

int main(int argc, char**argv) {
    printf("start num is %d\n", num);
    num++;
    int rc = fork();
    if(rc>0)
    {
        num+=2;
        printf("Parent's num is %d\n", num);
    }
    else if(rc==0)
    {
        num+=3;
        printf("Child's num is %d\n", num);
    }
    num+=4;
    printf("end num is %d\n", num);
}
```

Child's execution (blue)

```
int num = 3;

int main(int argc, char**argv) {
    printf("start num is %d\n", num);
    num++;
    int rc = fork();
    if(rc>0)
    {
        num+=2;
        printf("Parent's num is %d\n", num);
    }
    else if(rc==0)
    {
        num+=3;
        printf("Child's num is %d\n", num);
    }
    num+=4;
    printf("end num is %d\n", num);
}
```

Process Creation: *fork()*, and *wait()*

- Do you notice any difference? Yes, there is a *wait()* system call
- So what is the result?

```
prompt> ./p2
hello (pid:29266)
child (pid:29267)
parent of 29267 (rc_wait:29267) (pid:29266)
prompt>
```

```
int main(int argc, char *argv[]) {
    printf("hello (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {                // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("child (pid:%d)\n", (int) getpid());
    } else {                    // parent goes down this path
        int rc_wait = wait(NULL);
        printf("parent of %d (rc_wait:%d) (pid:%d)\n",
               rc, rc_wait, (int) getpid());
    }
    return 0;
}
```

- Variation: *waitpid*, *waitid*,

...

Process Creation: *fork()*, *wait()*, and *exec()*

- A final and important piece is the *exec()* system call.
 - Used when you want to run a program that is different from the calling program. (Unlike *fork()*, *exec()* will load a different program)
- The new process will replace the current process that is calling *exec()*.
 - This means statements **after** *exec()* call **will NOT be executed**
 - So don't put any statements after *exec()*.
- For example: calling *fork()* in previous example is only useful if you want to keep the running copies of the same program. However, often you want to run a different program --> *exec()* will do that.

Process Creation: *fork()*, *wait()*, and *exec()*

- In this example, the child process calls *execvp()* in order to run the program **wc**, which is the word counting program. In fact, it runs **wc** on the source file, thus telling us how many lines, words, and bytes are

```
prompt> ./p3
hello (pid:29383)
child (pid:29384)
      29      107      1030 p3.c
parent of 29384 (rc_wait:29384) (pid:29383)
prompt>
```

```
int main(int argc, char *argv[]) {
    printf("hello (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {                // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc"); // program: "wc"
        myargs[1] = strdup("p3.c"); // arg: input file
        myargs[2] = NULL;          // mark end of array
        execvp(myargs[0], myargs); // runs word count
        printf("this shouldn't print out");
    } else {                    // parent goes down this path
        int rc_wait = wait(NULL);
        printf("parent of %d (rc_wait:%d) (pid:%d)\n",
               rc, rc_wait, (int) getpid());
    }
    return 0;
}
```

Process Creation: *fork()*, *wait()*, and *exec()*

- It is confusing, is it? Why is it designed in this way? Why not use a single call to do the job?
- We need to duplicate the process and then let another process replace it. Is it a waste of time to copy paste?

Process Creation: *fork()*, *wait()*, and *exec()*

- Why separating *fork()* and *exec()*?
- It has benefits:
 - *Fork()* can be used independently: If we need multiple processes from the same program, then fork can do the job. (i.e. a web server can create a new process for each client)
 - Separating *fork()* and *exec()* allows additional control: Redirect output to a file
 - `wc p4.c > p4.output`

```
} else if (rc == 0) { // child: redirect standard output to a file
    close(STDOUT_FILENO);
    open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC); // Executing these two lines in child before exec
                                                    // allows us to redirect output.

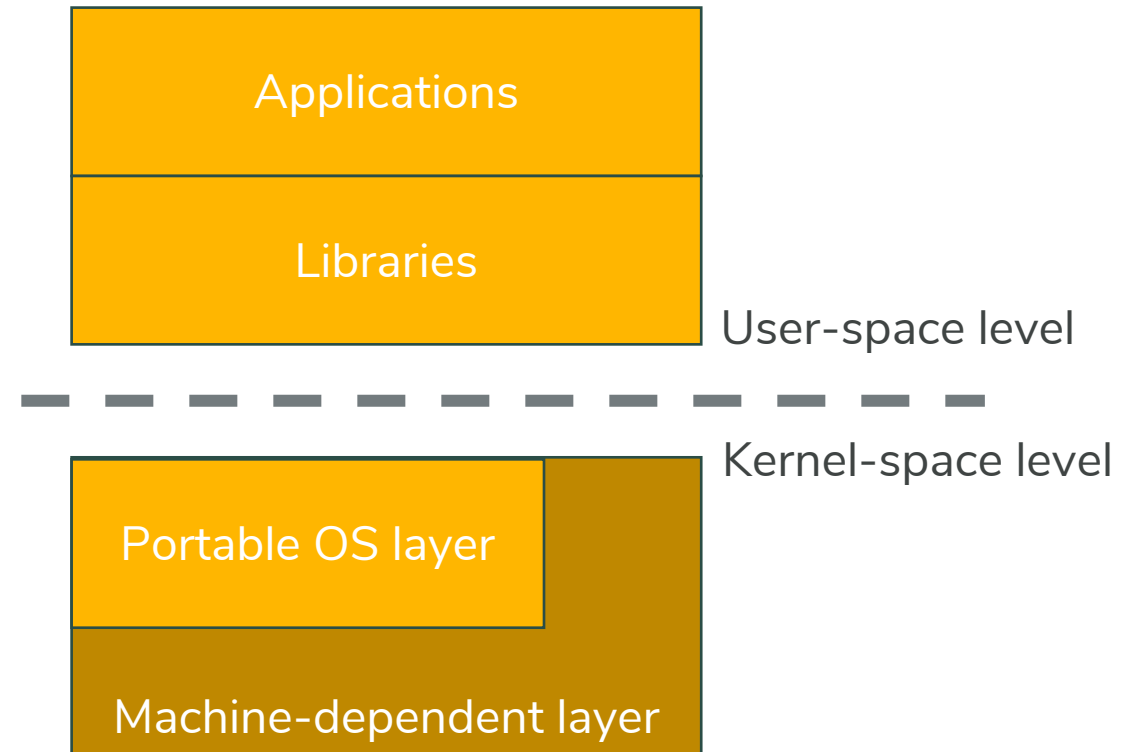
    // now exec "wc"...
    char *myargs[3];
    myargs[0] = strdup("wc"); // program: "wc" (word count)
    myargs[1] = strdup("p4.c"); // argument: file to count
    myargs[2] = NULL; // marks end of array
    execvp(myargs[0], myargs); // runs word count
} else { // parent goes down this path (main)
    int wc = wait(NULL);
}
```

Reduce copy overhead by copy-on-write

- Basic idea:
 - Don't copy memory when a process is forked
 - Only copy memory when either process is modifying memory
 - And only copy the part that is modified
- If calling `exec()` right after a `fork()`, almost nothing will be copied.
- We will learn more details later in memory management.


Process Memory Layout

- Remember this UNIX structure?
 - We need memory to store “user-space” and “kernel-space”
 - How can we arrange the storing?




Address Space

- One (common) approach:
 - Kernel is high memory
 - User is low memory
- Program segments:
 - Stack
 - Data
 - Text
 - Heap...




Address Space

- One (common) approach:
 - Kernel is high memory
 - User is low memory
- Program segments:
 - Stack
 - Data
 - Text
 - Heap...



Address Space

- Lots of flexibility
 - Allows stack growth
 - Allows heap growth
 - No pre-determined division.
 - Heap...



Cross-process communication (optional)

- Option 1: Share memory
- By default, process don't share memory
- System calls like mmap can create shared memory space among processes
- But be careful about all the correctness issues that can be caused by memory sharing: more details in Concurrency Chapter

Cross-process communication (optional)

- Option 2: Named Pipe
- What is a pipe? It's a UNIX mechanism that allows one process to write data to another process
- Named pipe: It has a well-known name (much like a filename) so that different processes can access it

Cross-process communication (optional)

- Option 3: Socket
- Key idea: Let processes communicate through networking
 - Socket is standard UNIX API to perform network communication
- Benefit: Processes can communicate even when they are not on the same machine

Cross-process communication (optional)

- Option 4: File
- How to know a file has been changed?
 - Periodic scanning is fine, but usually is not efficient.
 - In Linux, it provides system calls like inotify to monitor file changes.

Summary: Process (part 1)

- Process definition
 - Process States; PCB
 - Process Creation
 - Process Memory layout
 - Inter-process communication
-
- Next: Process Scheduling, Context Switch, and more on Inter-Process Communication