



# CSE2431 – Lecture Topic 2

## Process (part 2)







# Process (Part 2)

Instructor: Luan Duong, Ph.D.

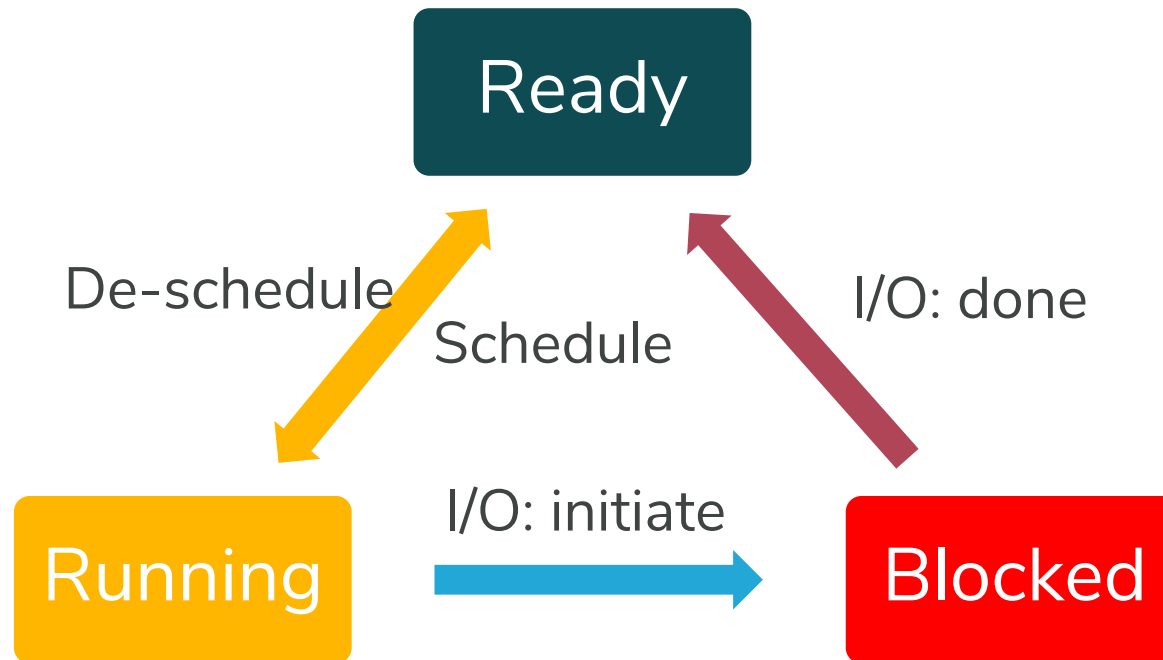
CSE 2431: Introduction to Operating Systems

**Reading: Chapter 4-5 in required textbook**

Lecture materials referred from previously taught course by  
Dr. Yang Wang and Dr. Adam C. Champion

# Last Lecture: Process status


- Running: a process is being executed by a CPU
- Ready: a process is ready to run but is not running
- Blocked: a process is waiting on some event to take place



# Last lecture: Process creation & memory

- *fork()*, *wait()* and *exec()*

```
int rc = fork();
if (rc < 0) {
    // fork failed
    fprintf(stderr, "fork failed\n");
    exit(1);
} else if (rc == 0) {
    // child (new process)
    printf("child (pid:%d)\n", (int) getpid());
} else {
    // parent goes down this path (main)
    printf("parent of %d (pid:%d)\n",
           rc, (int) getpid());
}
```

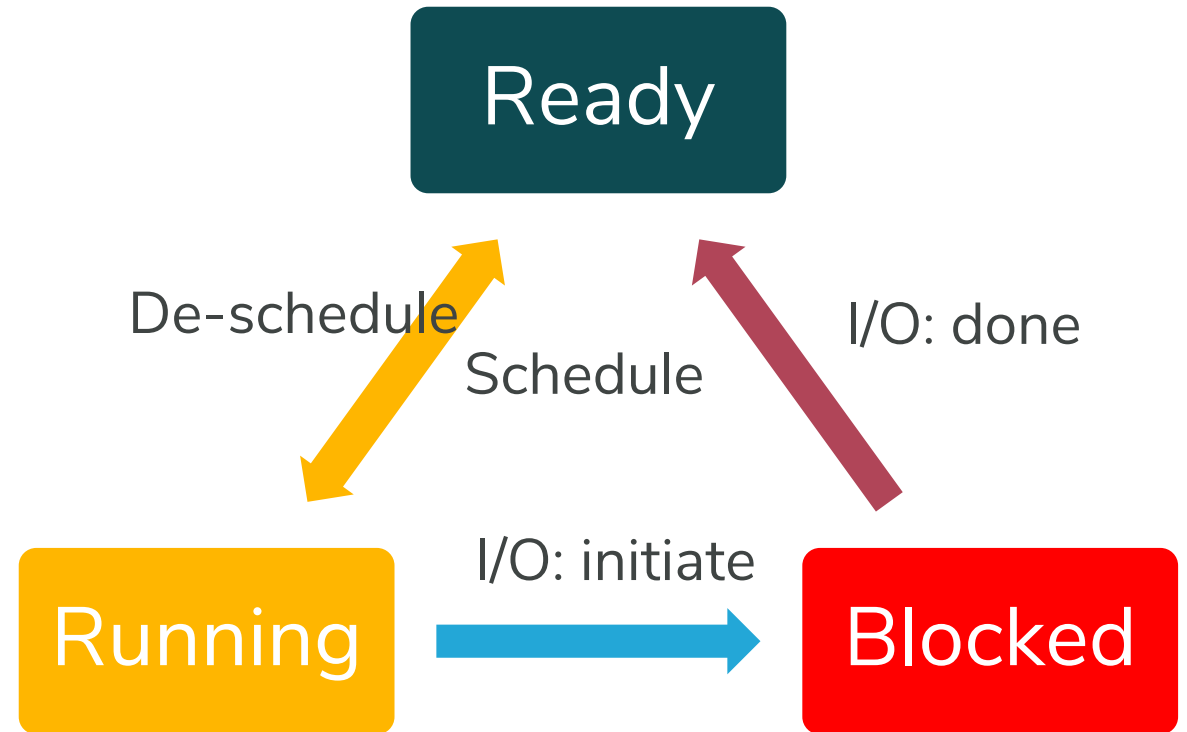


# Outline: Process

- What is a process?
- Process States; Process Control Block (PCB)
- Process Creation; *fork* command
- Process Memory Layout
- Process Scheduling
- Context Switch
- Inter-Process Communication
- Client-Server Communication

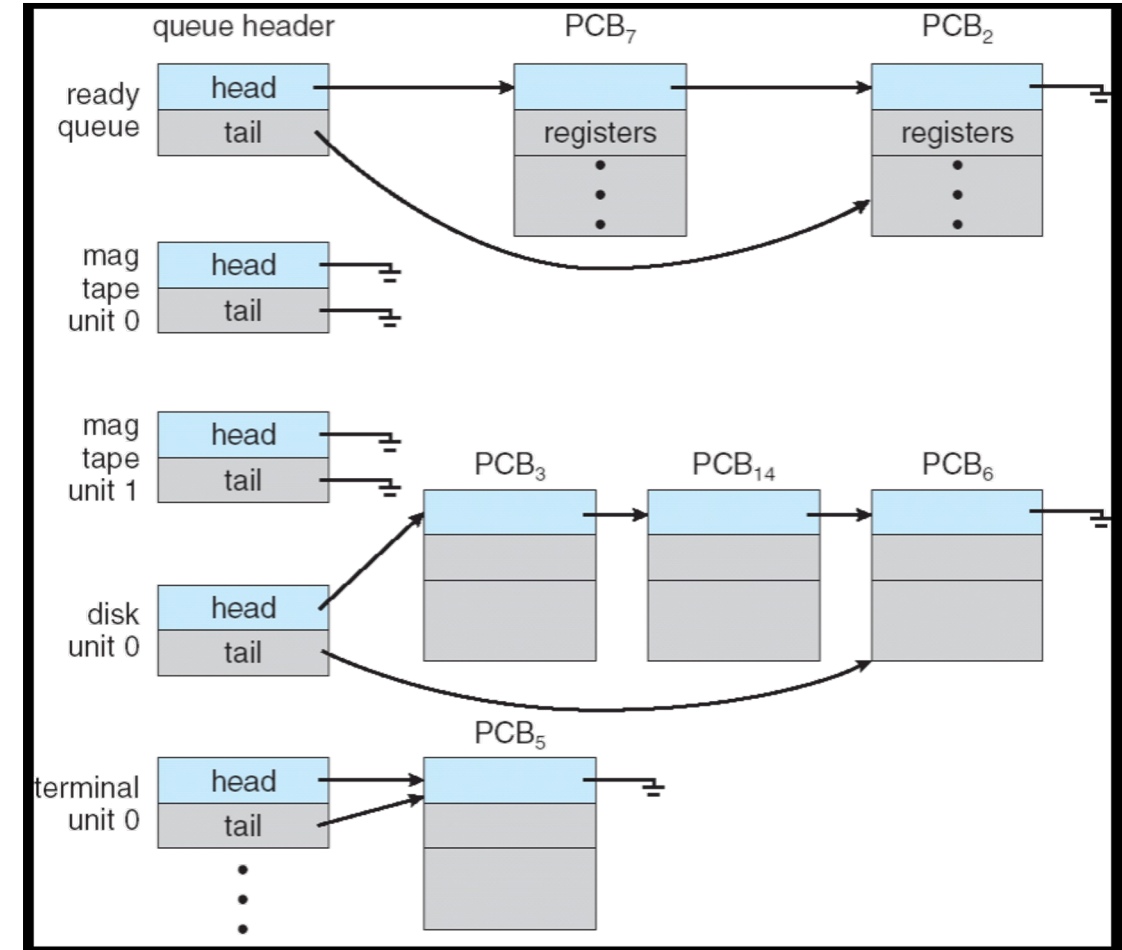
# Process scheduling

- Remember this?
  - A process is the basic unit for virtualization.
  - OS virtualizes CPU by **time sharing**
    - OS runs one process for some time, stops it, then runs another process, ...
    - Each process becomes slower, but users cannot tell, because computers are much faster than human beings
  - How can OS control the behavior of a process?



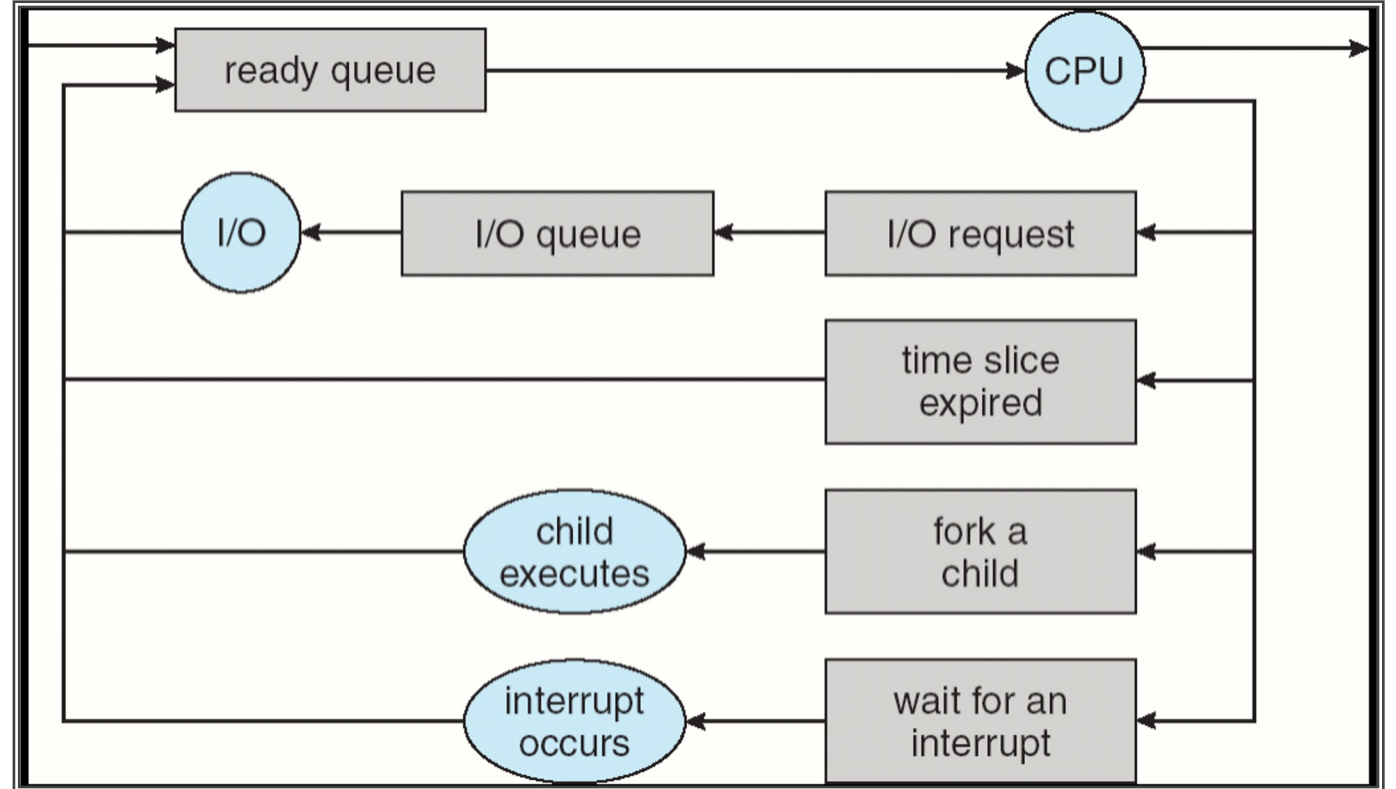
# Process scheduling

- **Job queue:** set of all processes in the system
- **Ready queue:** set of all processes in main memory, ready and waiting to run
- **Device queues:** set of processes waiting for an I/O device
- **Processes migrate among various queues**



# Process scheduling

- **Job queue:** set of all processes in the system
- **Ready queue:** set of all processes in main memory, ready and waiting to run
- **Device queues:** set of processes waiting for an I/O device
- **Processes migrate among various queues**



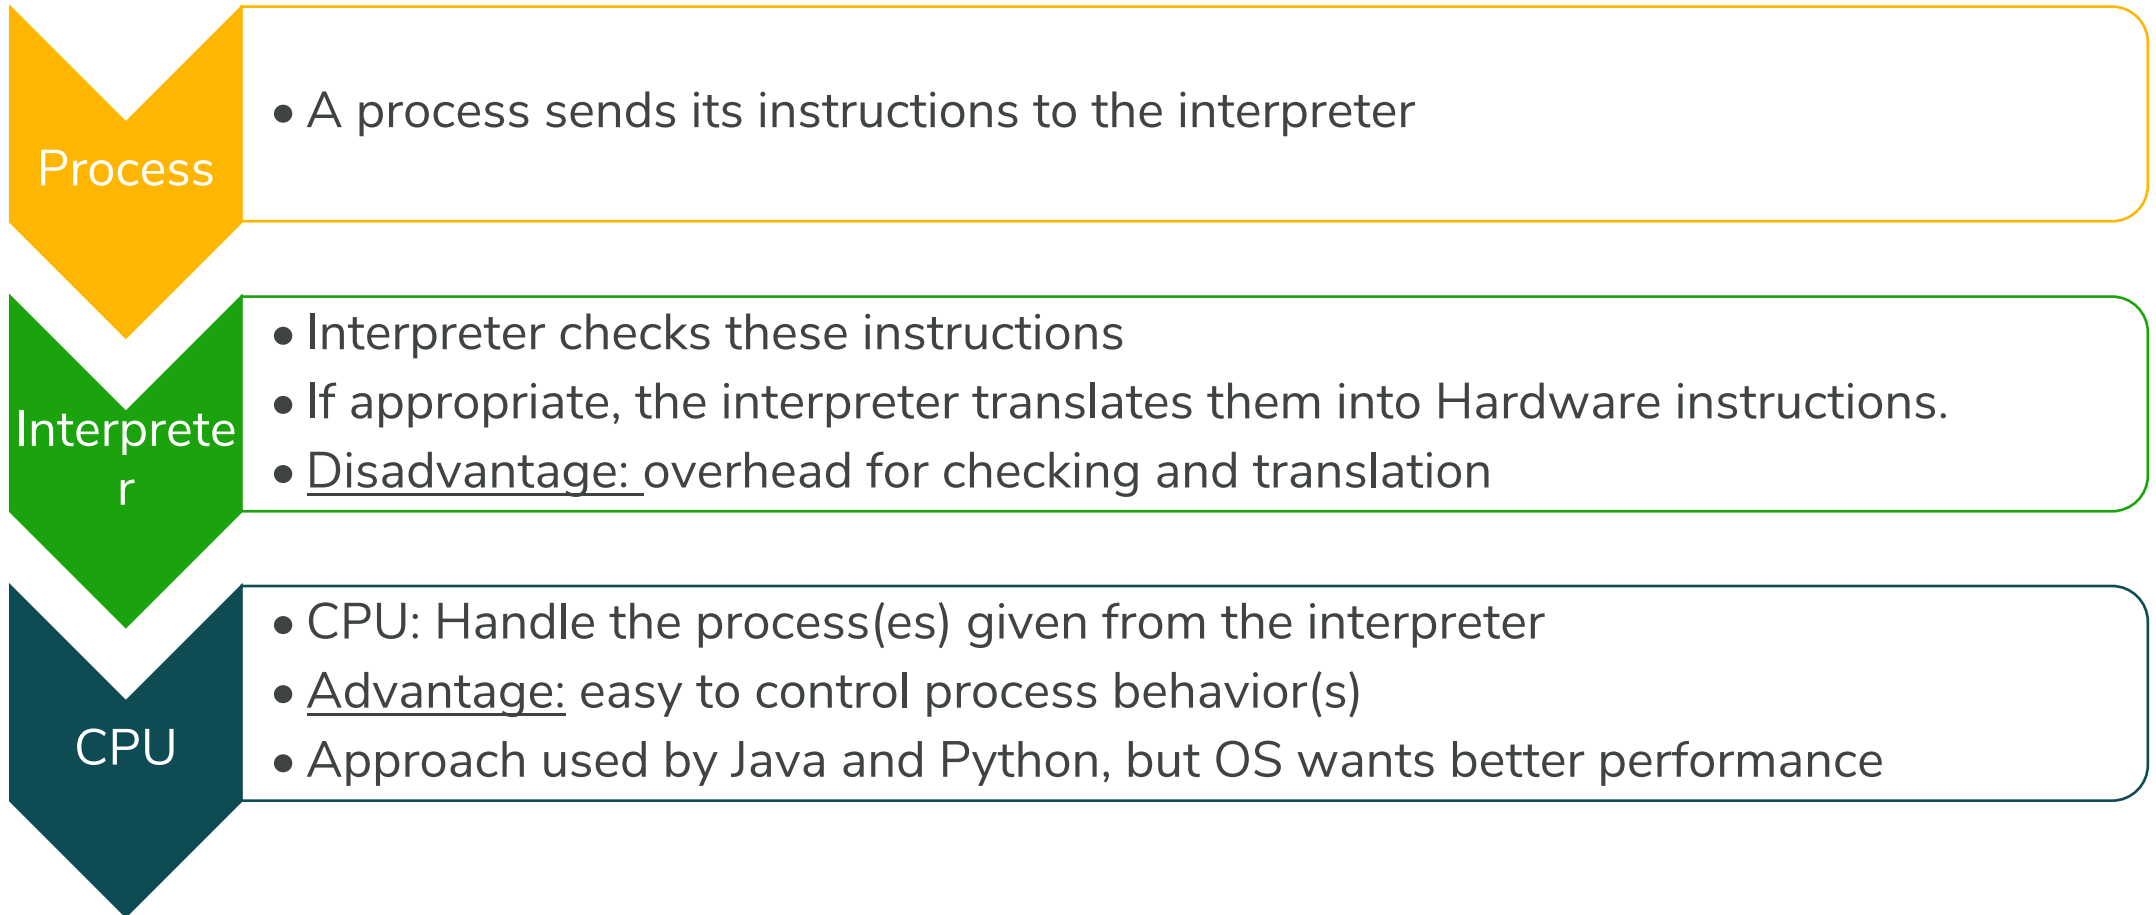


# Process scheduling

- **Long-term scheduler** (or job scheduler)
  - Selects which processes should be brought into the ready queue
  - Invoked very infrequently (seconds, minutes)  $\Rightarrow$  (may be slow)
  - Controls the degree of multiprogramming
  - Should balance different types of processes:
    - **I/O-bound process** spends more time doing I/O than computation; many short CPU bursts
    - **CPU-bound process** spends more time doing computation; few long CPU bursts
- **Short-term scheduler** (or CPU scheduler)
  - Selects which process should be executed next and allocates CPU
  - Invoked very frequently (milliseconds)  $\Rightarrow$  must be fast!


# Process scheduling

- Solution 1: Interpreter



# Process scheduling


- Solution 2: Limited Direct Execution
  - A process sends its instructions directly to the CPU
  - OS collaborates with CPU to control the behavior(s) of the process.
  - Advantage: low overhead
  - Disadvantages: hard to control process behavior(s)
  - OS usually uses this approach



# Process scheduling

- Solution 2: **Challenges**


- When a process starts to run, it gains **full control** of the CPU
- How to prevent a process from doing bad things?
  - A program can be malicious or simply buggy
  - It may try to read/write/delete other processes' data
  - This may violate virtualization
- How to pause a process and switch to another process?
  - A malicious or buggy program may not want to give up CPU resources
  - Ideally, we should not need to worry about this when writing a program





# Process scheduling

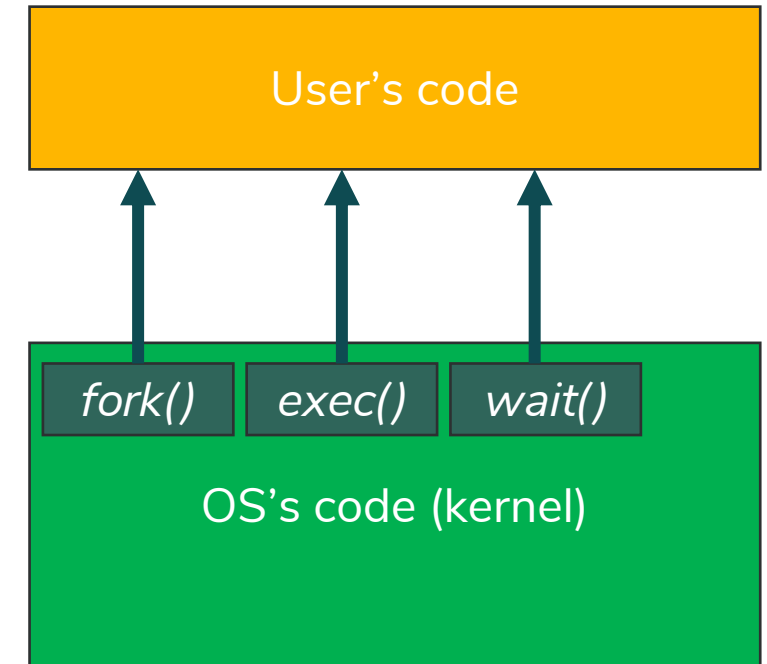
- **Solution?**
  - ‘Dangerous’ instructions will be called **restricted operations** or **privileged instructions**.
  - What instructions are not ‘dangerous’?
    - Normal operations such as arithmetic (add, sub, multiple)
    - Most memory-related operations (load, jump, conditional jump, etc.)
  - What are ‘dangerous’ instructions
    - i.e. write a file
- Solution: **system calls!**



# Process scheduling


- **Solution**

- User's code cannot execute restricted operations
- OS expose some functions to user's code
- OS code can execute restricted operations:
  - Define functionalities for users: fork/exec/wait; file create/open/read/write/close, ...
  - Execute restricted operations with appropriate checks (file write can only access data belonged to the user)
  - Programmer calls a system call like a library function.



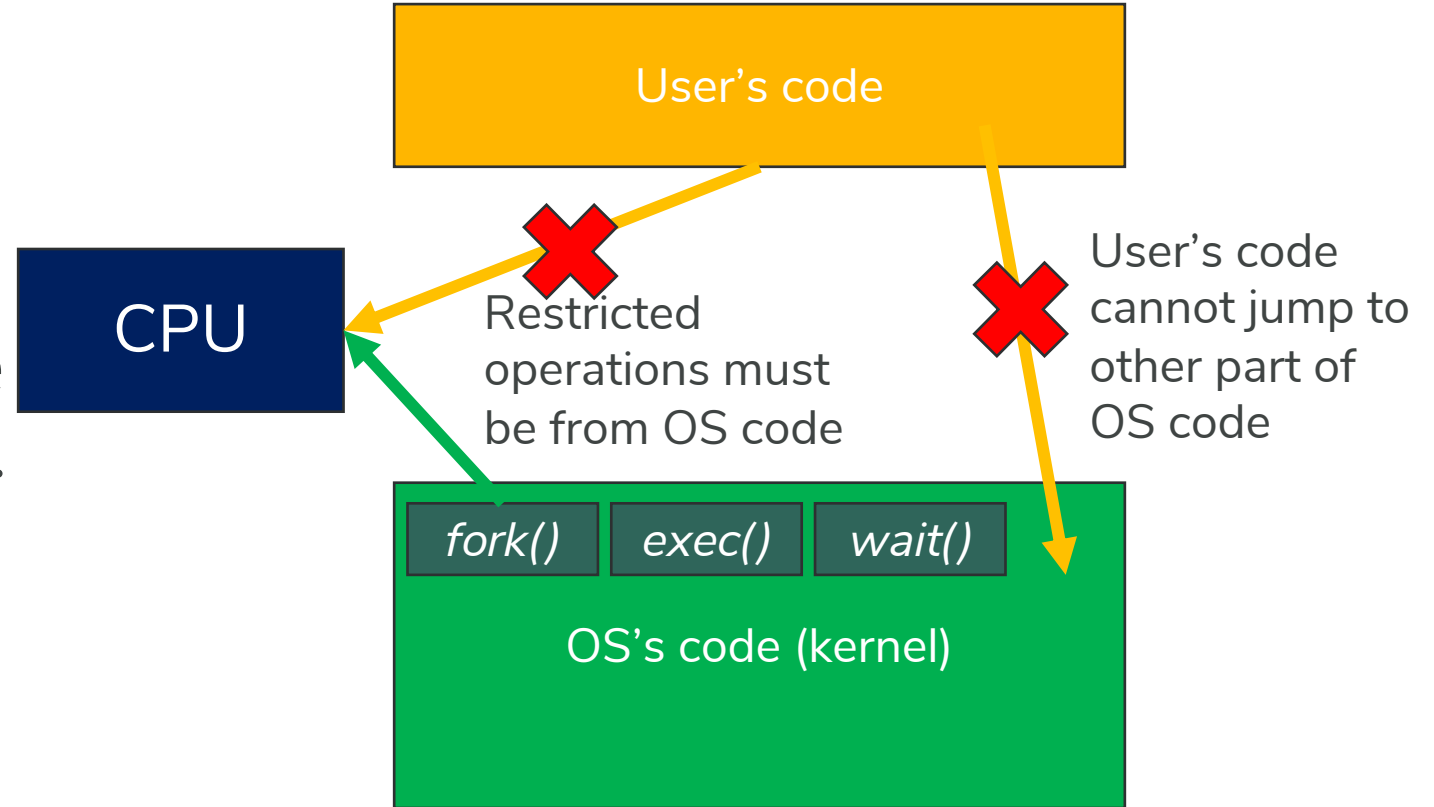
# Process scheduling

- Is the problem really solved?
  - **NO!**
  - We already prevent a user program from executing restricted operations arbitrarily.
  - However, we still assume a program is not buggy and not malicious.
  - A buggy/malicious program can still issue restricted operations without making system calls?
    - Remember this? How can CPU tell whether a restricted operation is from a system call or not? CPU only sees a sequence of instructions



# Process scheduling

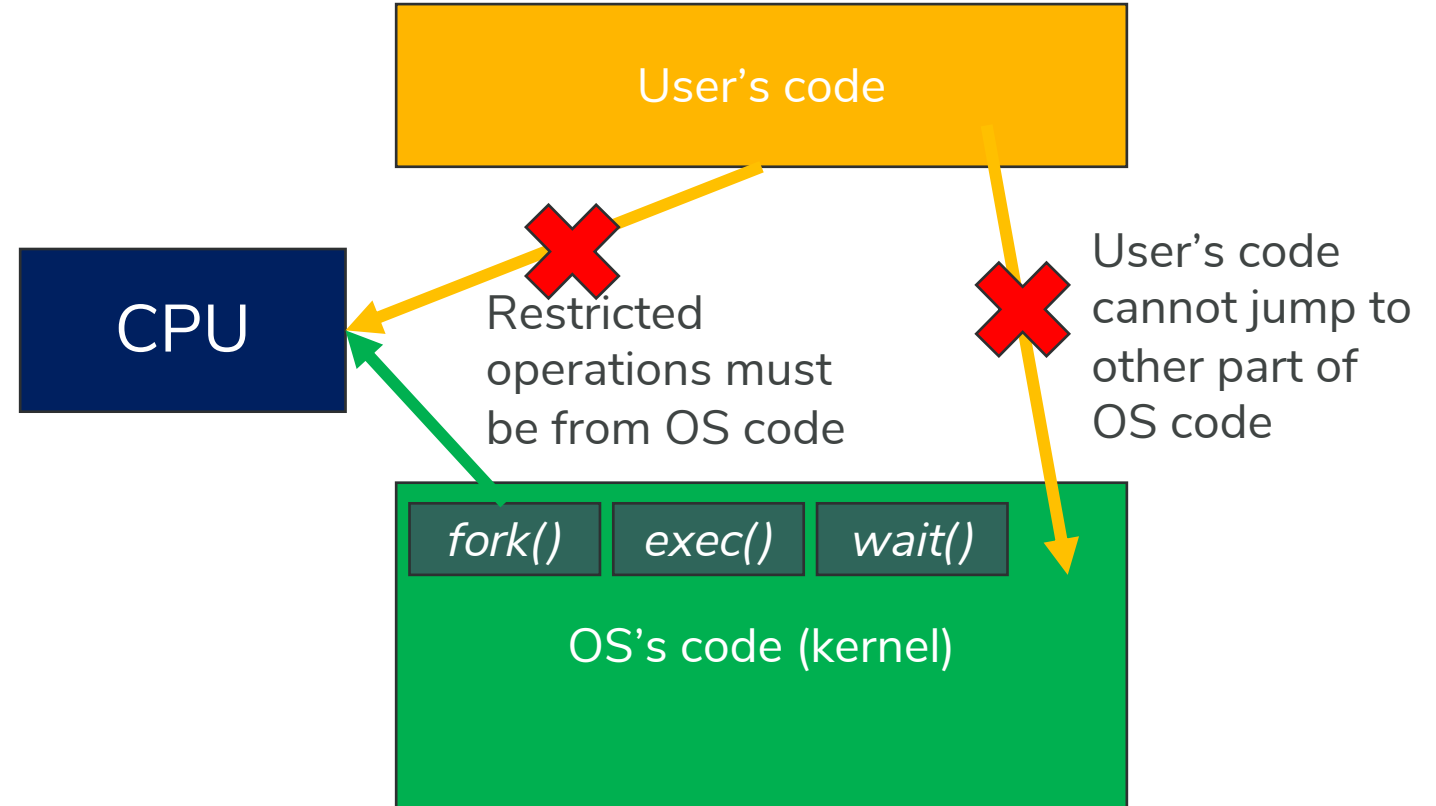
- Requirement?
  - User's code **cannot** execute **restricted operations**.
  - OS code **can** execute **restricted operations**.





# Process scheduling

- CPU's supports:
  - Trap-table
  - Kernel-mode and user-mode
  - Trap and return-from trap instruction



# Process scheduling

- Trap-table:
  - Normal function calls are implemented by “jump” instruction. Inside a process, code can jump to arbitrary location!
  - When calling an OS function, this **cannot happen** arbitrarily
  - OS exposes locations of all system calls in a **trap-table**
  - CPU ensures that user’s code can only jump to locations defined in trap-table.

# Process scheduling

- Kernel-mode and user-mode
  - To ensure only OS code can execute restricted operations
  - CPU maintains a bit
  - When this bit is set to True (**kernel mode**), OS code (including system call code) is executing
    - It can do anything including executing restricted operations.
  - When this bit is False (**user mode**), user code is executing
    - Only non-restricted operations are allowed
  - Have we really solved the problem with this?

# Process scheduling

- Kernel-mode and user-mode
  - Suppose CPU provides a new instruction “**change kernel bit**”
    - Should it be a restricted operation or not?
  - Can we make it **non-restricted**?
    - NO. A malicious code can set kernel bit to True and executed restricted operations.
  - Should we make it **restricted**?
    - Sure, but how can the CPU check that it is from OS code?
  - How can CPU knows whether the kernel bit itself is set up appropriately?

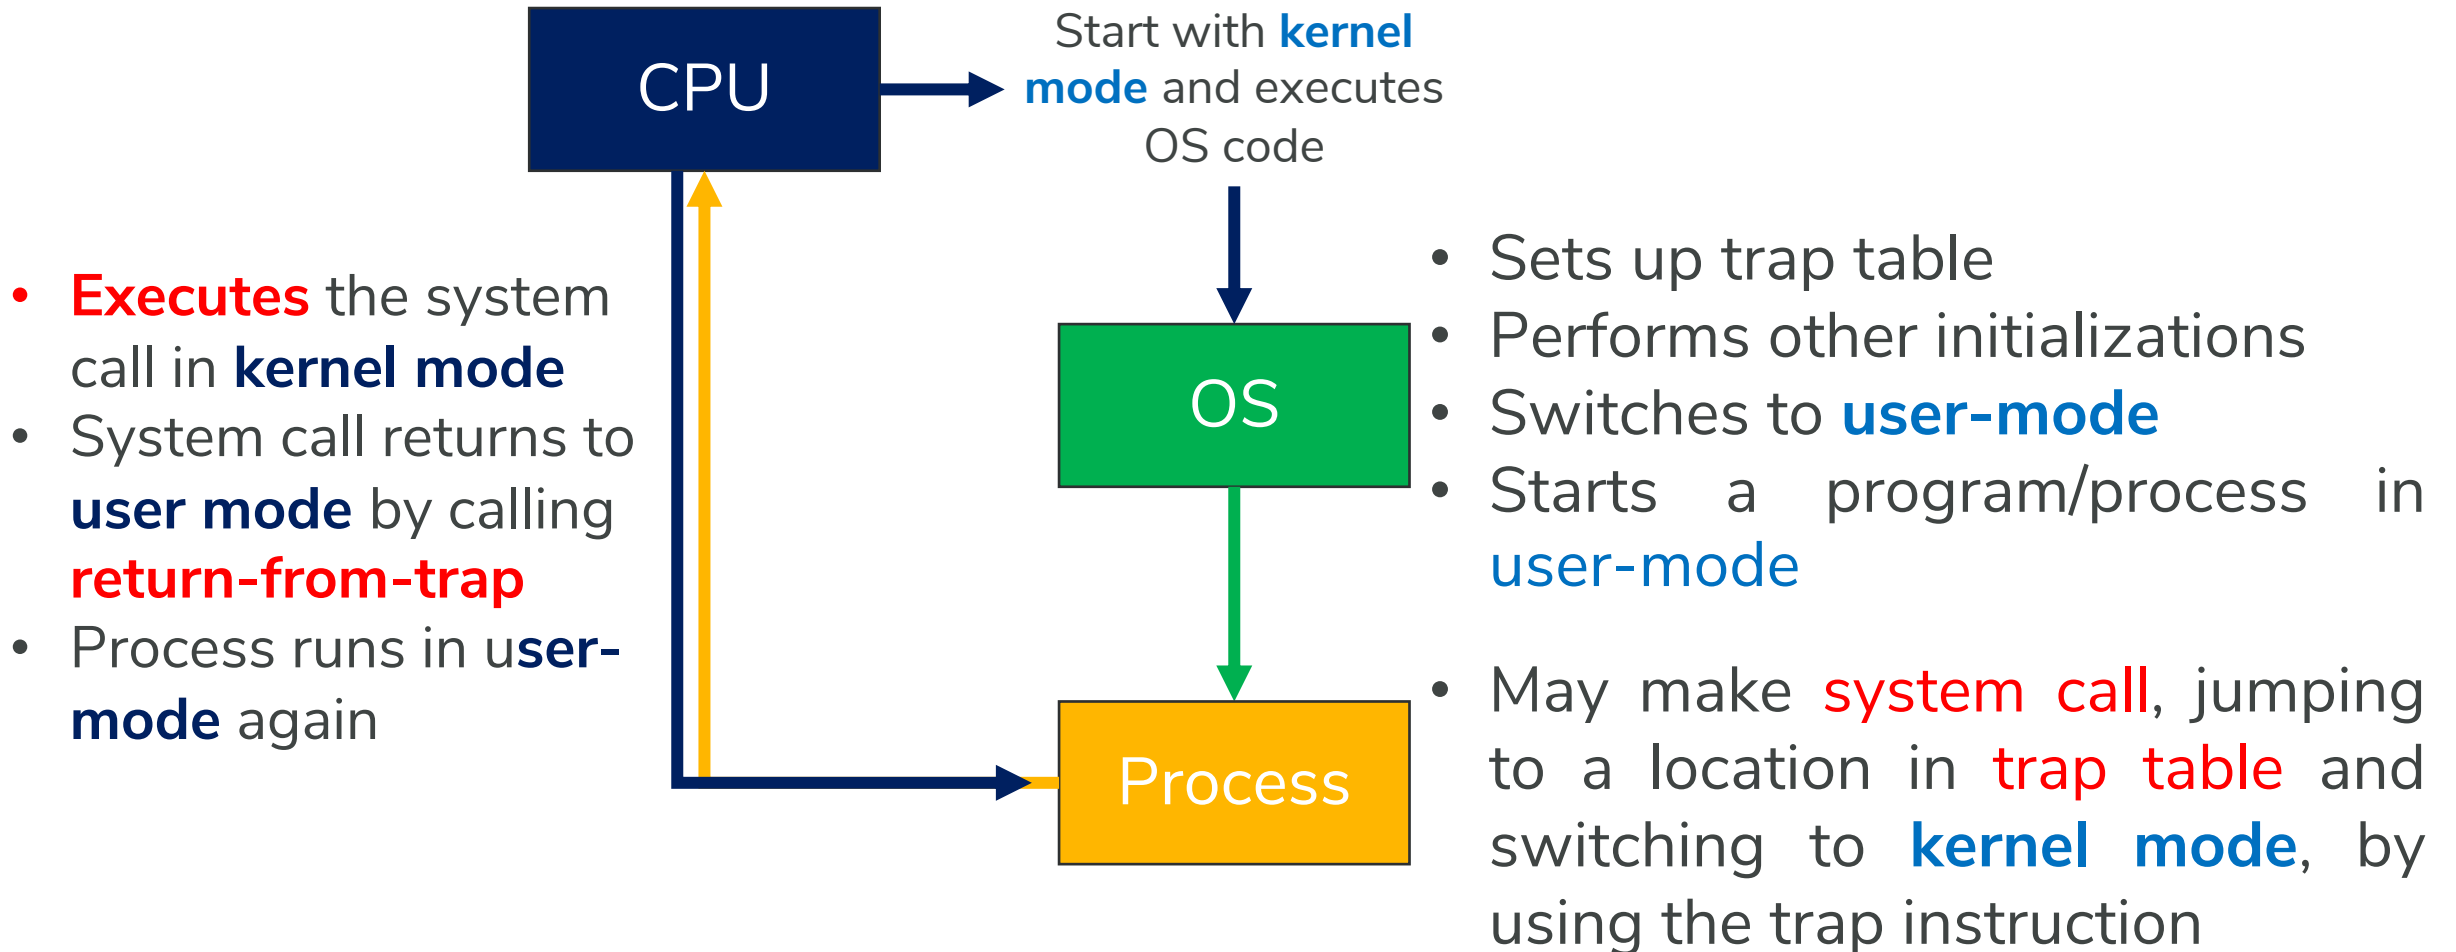


# Process scheduling

- Trap- and return-from-trap instruction
  - Modern CPU provides a special trap instruction that performs the following two operations **atomically**.
    - **Switch** from user mode to kernel mode
    - **Jump** to an address defined in the trap table
  - Key idea: we allow user's code to change kernel bit, but **after** doing so, the user's code **must jump** to a **system call**!
    - Then it is safe to make trap non-restricted.
  - **OS** makes sure a system call **switches to user mode** before returning
    - CPU provides a special return-from-trap instruction to do that

# Process scheduling

- At boot time:



# System call vs. normal function call

System call	Function call
Must run in <b>kernel mode</b>	Can run in any mode
Made by a <b>trap instruction</b>	Made by a jump instruction
Must be registered in the <b>trap table</b>	Can be put in any location (in mem)
Comes with <b>two context switches</b>	Only with one context switch

# Outline: Process

- What is a process?
- Process States; Process Control Block (PCB)
- Process Creation; *fork* command
- Process Memory Layout
- Process Scheduling
- **Context Switch**
- Inter-Process Communication
- Client-Server Communication



# How to pause a process?

- Remember when a process is running, OS is not running
- Cooperative approach: OS trusts the process to behave reasonably
  - OS can take control when the process is making a system call
  - OS provides an explicit “yield” system call
  - Problem: it violates virtualization; malicious or buggy processes can still halt the whole system by not making any system calls
- Non-cooperative approach: OS takes control anyway
  - Need hardware support

# How to pause a process?

- Special hardware support—timer
  - A timer will trigger a timer event periodically
  - When a timer event is triggered, CPU will automatically jump to a predefined timer event handler
  - In this way, CPU provides a backdoor for OS.
- An OS can implement its own timer event handler and register it to the CPU (much like registering a system call)
  - An OS can implement the process switching logic there

# How to resume a process?

- Remember a CPU decides the next instruction to execute by the program counter (PC) register
  - So, to resume a process, an OS should set PC to the location where the process is paused
- There are some other registers the OS needs to recover.
- In principle, an OS should save the state of a process when pausing it and load the state when resuming the process.
  - This is called a **context switch**.
  - It's very much like saving, loading PC game states

# Context Switching

- Switches CPU from one process to another
- Performed by scheduler (dispatcher) with a special hardware support (timer)
- It includes:
  - Saving the state of the old process (such as registers);
  - Loading the state of the new process;
  - Flushing memory cache;
  - Changing memory mapping [Translation Lookaside Buffer (TLB)]

# Context Switch Workflow

## ① Process 1

```
0x00: load r1, &a
0x04: inc r1
0x08: store r1, &a
0x0C: load r1, &a
0x10: cmp r1, 1
0x14: JE 0x20
0x18: add r1, 2
0x1C: store r1, &a
0x20: .....
```

CPU

PC=0x08

## ② Timer event occurs

## ③ Save PC=0x08



# Context Switch Workflow

## ① Process 1

```
0x00: load r1, &a
0x04: inc r1
0x08: store r1, &a
0x0C: load r1, &a
0x10: cmp r1, 1
0x14: JE 0x20
0x18: add r1, 2
0x1C: store r1, &a
0x20: .....
```

## ③ Save PC=0x08

CPU

PC=0x08

## ② Timer event occurs

## Process 2

```
0x50: .....
0x54: .....
0x58: .....
0x5C: .....
0x60: .....
0x64: .....
0x68: .....
0x6C: .....
0x60: .....
```

## ④ Timer event handler

# Context Switch Workflow

## ① Process 1

```
0x00: load r1, &a
0x04: inc r1
0x08: store r1, &a
0x0C: load r1, &a
0x10: cmp r1, 1
0x14: JE 0x20
0x18: add r1, 2
0x1C: store r1, &a
0x20: .....
```

## ② Save PC=0x08

CPU

PC=0x08

- ③ Timer event occurs
- ④ Jump to timer-event handler
- ⑥ OS resumes Process 1 **or** may switch to Process 2

## Process 2

```
0x50: .....
0x54: .....
0x58: .....
0x5C: .....
0x60: .....
0x64: .....
0x68: .....
0x6C: .....
0x60: .....
```

## ⑤ Timer event handler

# Context Switch (more)

- Context switches can occur when
  - A program makes a system call
  - A timer event is triggered
  - .....
- Context switches have additional overhead
  - OS needs to save the state of current process and load the state of another process
  - Frequent context switches can be bad for performance



# Context Switching

- **Context switch is expensive** (1–1000  $\mu$ sec)
  - No useful work is done (pure overhead)
  - Can become a bottleneck
  - **Real-life analogy?**

# Summary: Process (part 2)

- Process Scheduling
- Context Switching
- Inter-Process Communication