

I/O and File System (Part 1)

Instructor: Luan Duong, Ph.D.

CSE 2431: Introduction to Operating

Systems

Reading: Chap. 36, 37, and 38 [OSTEP]



Outline: I/O and File System (part 1)

- I/O Systems
 - What are I/O Devices?
 - Disks (again):
 - Disk Performance
 - Disk Scheduling
 - I/O System
 - I/O System Architecture
 - I/O Devices Operating



What is I/O device?

- Apart from CPU and memory, almost everything else is an I/O Device.
- Examples: disk, graphic card, screen, mouse, keyboard, network card, printer, joystick, sound card, ...

- How does an OS operate so many kinds of I/O devices?
 - To answer this question, let's first see how I/O devices work.



Disk Performance: Speed

- Disks are mechanical, so they're slow compared to CPUs
 - Seek: several msec (usually)
 - Rotation: typical/common speeds include 5400 RPM (rotations per minute), 7200 RPM, and 15000 RPM.
 - Average rotation time (seconds) = 1/ (RPM/60) / 2 (several msec)
 - Data transfer: usually 100 MB/sec – 150 MB/sec

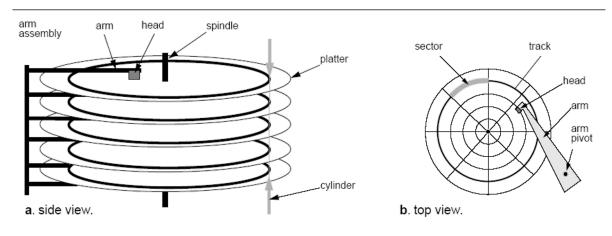


Figure 1: the mechanical components of a disk drive.



Disk Performance: Access Time (1)

Average access time for a sector:

$$T_{\text{access}} = T_{\text{avg_seek}} + T_{\text{avg_rotation}} + T_{\text{avg_transfer}}$$

Seek time (**T**_{avg_seek}):

- Time to position heads over cylinder
- Typical $T_{\text{avg seek}}$ is 3–9 msec, 20-msec maximum

Rotational latency (*T*_{avg_rotation}):

- After head is positioned over track, the time it takes for the first bit of the sector to pass under the head
- Worst case: head just misses the sector and waits for the disk to rotate 360

$$T_{\text{max_rotation}} = (1/\text{RPM}) \times (60 \text{ secs/1 min})$$

Average case is half of worst case:

$$T_{avg\ rotation} = (1/2) \times (1/RPM) \times (60 \text{ secs/1 min})$$

• Typical $T_{avg_rotation} = 5400-7200 \text{ RPM}$



Access Time (2)

Transfer time (T_{avg_transfer}):

- Time to read bits in the sector
- Time depends on rotational speed, number of sectors per track.
- Estimate of the average transfer time:
 - $T_{\text{avg_transfer}} = (1/\text{RPM}) \times (1/(\text{avg #sectors/track})) \times (60 \text{ secs/1 min}) \times (1000 \text{ msec/1 sec})$

Example:

- Rotational rate = 7200 RPM
- Average seek time = 9 msec
- Avg #sectors/track = 400

```
T_{\text{avg\_rotation}} = 1/2 \times (60 \text{ secs/}7200 \text{ RPM}) \times (1000 \text{ msec/sec}) = 4 \text{ msec}
T_{\text{avg\_transfer}} = (60/7200 \text{ RPM}) \times (1/400 \text{ secs/track}) \times (1000 \text{ msec/sec}) = 0.02 \text{ msec}
T_{\text{access}} = 9 \text{ msec} + 4 \text{ msec} + 0.02 \text{ msec} = 13.02 \text{ msec}
```



Exercise

- Consider the Cheetah 15K.5 disk with these parameters:
 - 15,000 RPM
 - 4-msec average seek time
 - Max disk transfer: 125 MB/sec
- How long does it take to transfer 4 KB of data (assuming data, sectors are aligned)?
- How long does it take to transfer 4 MB of data?



Exercise: Solved

- Seek time = 4 msec
- Rotation time = 1/(15000/60)/2 = 2 msec
 - 1. For 4 KB of data: transfer time = 4 KB / (125 MB/sec) = 0.031 msec
 - Transfer time is almost negligible compared to seek and rotation time
 - Disk access time = 4 msec + 2 msec + 0.031 msec = 6.031 msec
 - 2. For 4 MB data:, transfer time = 4 MB / (125 MB/sec) = 31 msec
 - Disk access time = 4 msec + 2 msec + 31 msec = 37 msec
 - Compared to case (1), we transfer $1,000 \times$ the amount of data in only $6 \times$ the amount of time!
- Disks prefer sequential, large I/O operations



Disks Prefer Sequential, Large I/O operations

- Sequential, large I/O operations can be 100× faster than small, random ones
 - OS has mechanisms to assemble small, sequential I/O operations into a large I/O one, so these two are somewhat equivalent.
 - The disk always spins!
 - Hence, if you issue an I/O operation to sector 1, wait awhile, and then issue another I/O operation to sector 2, they're *not* considered sequential I/O operations.
 - The disk and OS determine the acceptable interval between these.
- Reason: seek and rotation times dominate.
 - New devices like solid-state drives (SSDs) and non-volatile memory (NVM) have alleviated this problem
 - Yet, even on these devices, sequential I/O is much faster than random I/O



Unaligned I/O

- Thus far, we've assumed that I/O is aligned with sectors
- What do we do to read or write half a sector or 1.5 sectors?
 - Recall: a disk's minimal I/O data unit is a sector
- Read is simpler: just read all corresponding sectors that contain our targets, and then copy target data out
- Write: we must read all sectors first, modify them, then write back
 - This is called a read-modify-write procedure
 - It further slows down a disk, because a write needs an extra read



Lesson to Learn

- Avoid random I/O operations
- Avoid unaligned I/Os, in particular unaligned writes!
- These optimizations significantly improve the performance of your programs



Disk Performance Factor: Seeking

- Seeking: position the head to the desired cylinder (2–5 msec)
- Seek speed depends on:
 - Available power for pivot motor $(0.5 \times \text{ seek time needs } 4 \times \text{ power})$
 - Arm stiffness (30–40g acceleration required for short seek time; flexible arms can twist ⇒ crash head into platter surface!)
- A seek is composed of
 - A speedup, a coast, a slowdown, a settle
 - For very short seeks, the settle time dominates (1–3 msec)
- Real-life analogy?



Questions

 Since seek time and rotation delay dominate, how do you improve I/O performance?

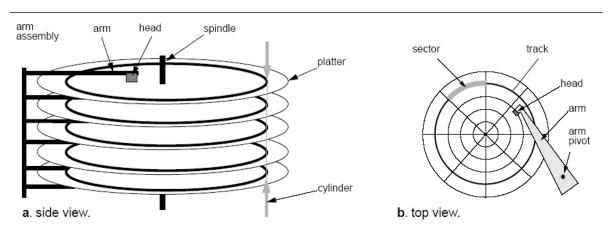


Figure 1: the mechanical components of a disk drive.



Outline: I/O and File System (part 1)

- I/O Systems
 - What are I/O Devices?
 - Disks (again):
 - Disk Performance
 - Disk Scheduling (Hard Disk Drive: HDD)
 - I/O System
 - I/O System Architecture
 - I/O Devices Operating



Hard Disk Scheduling

- Which disk request is serviced first?
 - FCFS
 - Shortest seek time first
 - Elevator (SCAN)
 - C-SCAN (Circular SCAN)
- Look familiar?



Hard Disk Scheduling: FIFO (FCFS)

Method

First come first serve

Pros

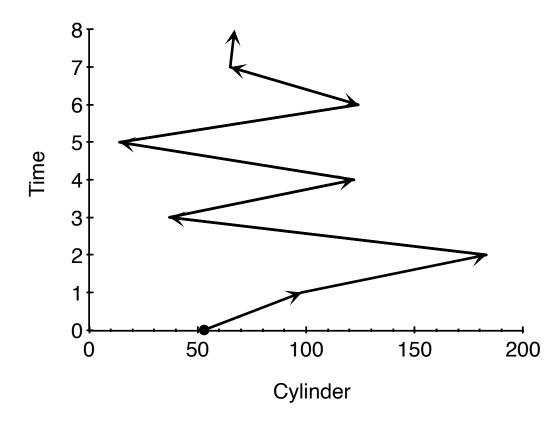
- Fairness among requests
- In the order applications expect

Cons

- Arrival may be on random spots on the disk (long seeks)
- Wild swing can happen

Analogy:

 Can elevator scheduling use FCFS?



Cylinders: 98, 183, 37, 122, 14, 124, 65, 67

Start position: cylinder 53

Direction: heading toward cylinder 200

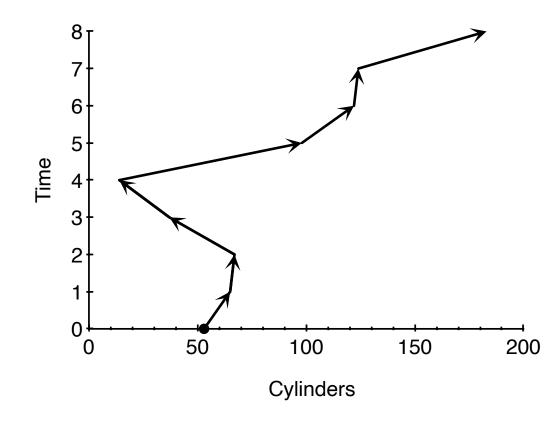
Traversal order: 53, 98, 193, 37, 122, 14, 124, 65, 67



SSTF (Shortest Seek Time First)

Method

- Pick the one closest on disk
- Rotational Delay is in calculation
- Pros
 - Aims to minimize seek time
- Cons
 - Starvation
- Analyses:
 - Is SSTF optimal?
 - Can we avoid starvation?



Cylinders: 98, 183, 37, 122, 14, 124, 65, 67

Start position: cylinder 53

Direction: heading toward cylinder 200

Traversal order: 53, 65, 67, 37, 14, 98, 122, 124, 183



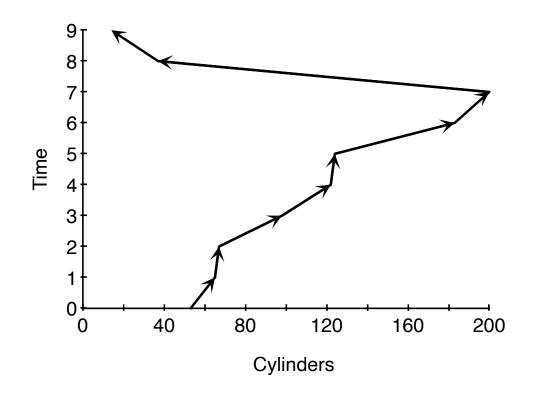
Elevator (SCAN)

Method

- Take the closest request in the direction of travel
- Travel to end, change direction
- Pros: Bounded time for each request
- Cons: Request at the other end will take time

LOOK algorithm

- Do not go to the end
- Service last request, then change direction



Cylinders: 98, 183, 37, 122, 14, 124, 65, 67

Start position: cylinder 53

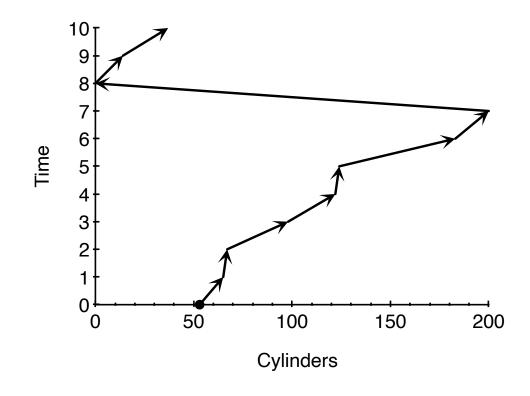
Direction: heading toward cylinder 200

Traversal order: 53, 65, 67, 98, 122, 124, 183, 200, 37, 14



C-SCAN (Circular SCAN)

- Method: like SCAN, but wrap around
- Pros: Uniform service time
- Cons: Does nothing on return
- C-LOOK
 - Do not go to the end
 - Service the last request, then wrap around



Cylinders: 98, 183, 37, 122, 14, 124, 65, 67

Start position: cylinder 53

Direction: heading toward cylinder 200

Traversal order: 53, 65, 67, 98, 122, 124, 183, 200, 0, 14, 37



SPTF: Shortest Positioning Time First

- So far, we've only considered track (seek time), but rotation also takes time.
- Accessing a sector with a farther track could be faster, due to less required rotation.
- Computing both track and rotation is too hard for OS
 - Hard disk devices often use SPTF internally



Other Issues

- Where to schedule disk I/O operations?
 - Both disks and the OS do that.
- I/O merging: OS can merge small I/Os into big ones
- How long should an OS wait before issuing an I/O?
 - Longer wait times may yield more merges, greater scheduling efficiency
 - But if no new I/O operations arrive, this wastes time!



Solid-State Disk Scheduling

- SSD scheduling algorithms simpler than HDDs
- Linux uses NOOP (no scheduling), but nearby logical block requests combined
 - Write amplification (one write leads to garbage collection, many I/O ops.)
 hurts performance
 - File system can notify about empty blocks to be erased (TRIM on SATA SSDs)
- More info: J. Kim, Y. Oh, E. Kim, J. Choi, D. Lee, and S. H. Noh, "Disk Schedulers for Solid State Drivers," in *Proc. ACM EMSOFT*, 2009.



History of Disk-related Concerns

- When memory was expensive, minimize bookkeeping
- When disks were expensive, maximize number of usable sectors
- When disks became more common, increase reliability
- When processors got faster, make them appear faster
- Key point: find, improve bottlenecks, single points of failure

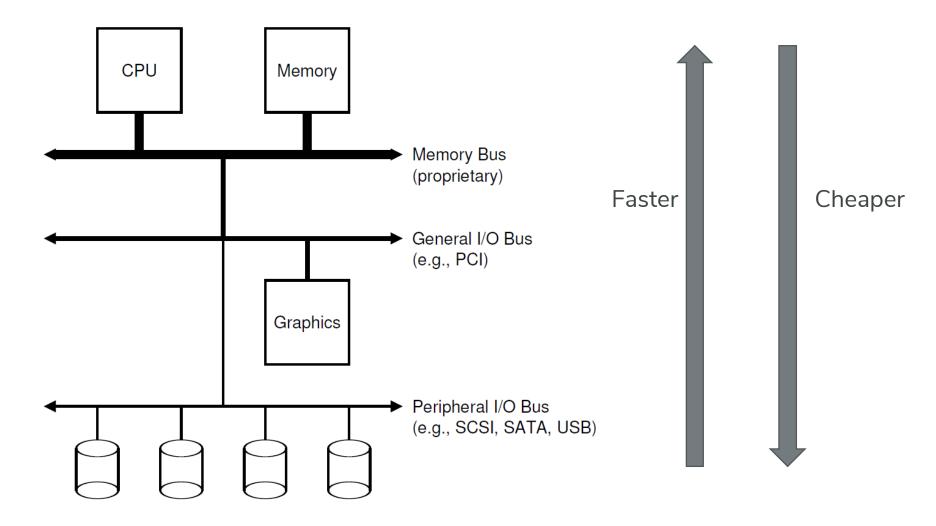


Outline: I/O and File System (part 1)

- I/O Systems
 - What are I/O Devices?
 - Disks (again):
 - Disk Performance
 - Disk Scheduling
 - I/O System
 - I/O System Architecture
 - I/O Devices Operating

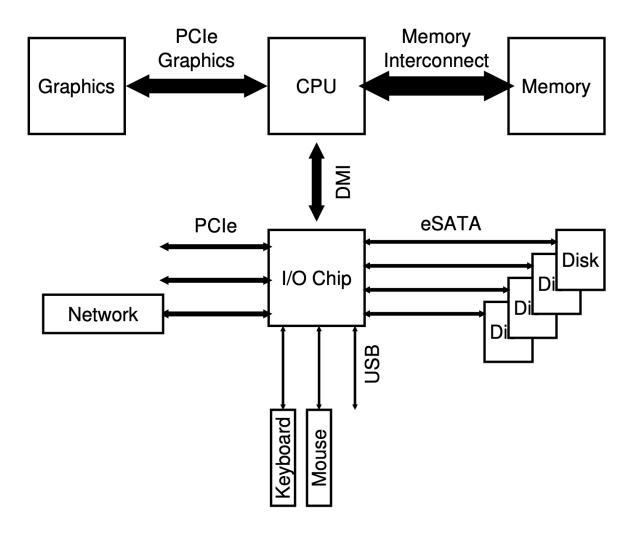


Architecture





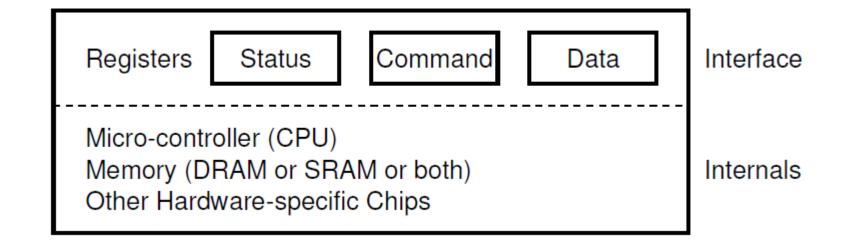
Architecture (Modern)





A canonical device

- Interface: usually a set of registers
- Internal structure: vary a lot





Operating the device

A typical piece of pseudo code

```
While (STATUS == BUSY)
   ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
   (Doing so starts the device and executes the command)
While (STATUS == BUSY)
   ; // wait until device is done with your request
```

- It continuously polls registers to check state (polling mode)
- It requires CPU to write data to registers (PIO mode)



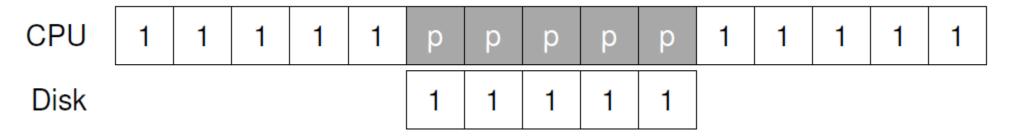
Polling vs. Interrupts

- Polling consumes CPU, so polling for long is not efficient
- Alternative solution: interrupt
 - OS registers an interrupt handler to CPU
 - Instead of polling, OS puts the process that is performing I/O to sleep
 - OS may switch to another process
 - When the I/O device is ready or has completed its job, it issues an interrupt
 - CPU calls interrupt handler and may wake up the original process

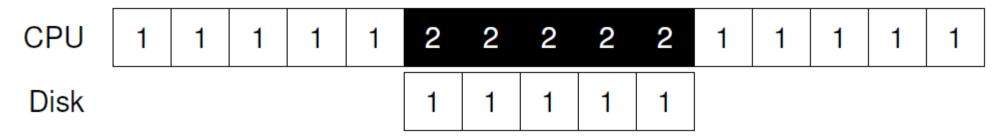


Polling vs. Interrupts

Polling mode



Interrupt mode





Polling vs. Interrupts

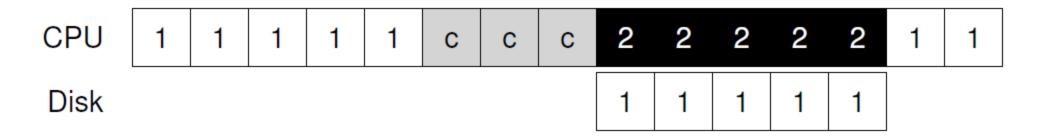
 Interrupt has its own overhead: it needs to perform two context switches

- For fast devices, polling may be a better solution. Why?
 - If every polling takes short, its overhead may be smaller than two context switches



Programmed I/O vs Direct Memory Access

- PIO (programmed I/O): CPU moves data from its own register to I/O devices' register (may need to move data from memory first)
 - CPU cannot do other things when moving data



• For fast I/O devices, PIO can consume a lot of CPU.



Programmed I/O vs Direct Memory Access

 Alternative solution: DMA (Direct Memory Access). Modern computer has some special hardware (DMA engine) that can transfer data across devices

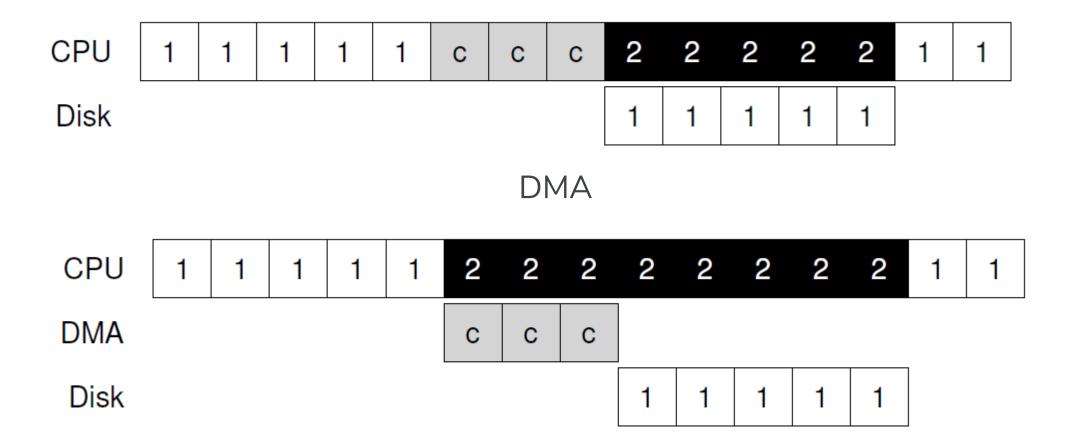
DMA usage:

- OS tells DMA engine location of data, size of data, and destination of data
- OS is done and can switch to other work; DMA engine copies data in the meantime
- When DMA is complete, DMA engine raises an interrupt
- OS realizes data copy is complete



Programmed I/O vs Direct Memory Access

PIO





How to read/write device registers?

- One solution: I/O instructions (e.g. "in" and "out" on x86)
 - They are privileged instructions
- Another solution: memory-mapped I/O
 - Map device registers to some specific memory locations
 - Then operate them with "load" and "store" instructions
 - These memory locations are not visible to user programs.
- No fundamental differences between these two solutions



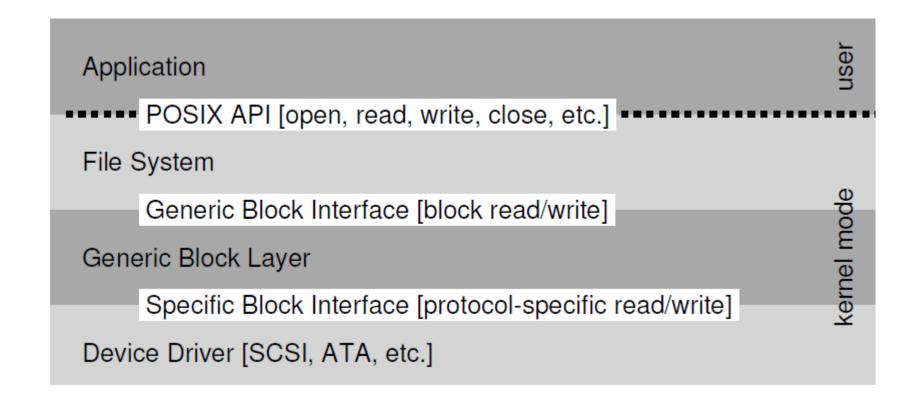
Go back to our original question

- How can OS operate so many kinds of I/O devices?
 - Each device has different registers, commands, and data formats.
 - New devices may be released after OS.
- After you learn to drive, you can drive different kinds of cars. How does this happen? Cars may have different engines, transmissions, etc.
 - Reason: all cars follow a similar specification (steering wheel, brake, etc).
- OS uses a similar solution
 - OS defines a specification for each type of I/O device (e.g. disk must implement read/write operations)
 - Device manufacturer implements this specification in a device driver (e.g. how to do read/write by operating disk registers)
 - After you install a new device, you must also install its device driver.



Device drivers

Multi-level abstractions





Facts about Device Drivers

- 70% of Linux kernel code is actually for device drivers
 - Windows probably has a similar number
- Compared to OS core code, device drivers are often not written by kernel programming experts and are often not well tested
- As a result, they are a primary contributor to bugs and kernel crashes
 - It may be a mistake to blame Microsoft for "blue screen"



Case study: A simple IDE disk driver

- High-level abstraction: OS requires all block devices (disk is a kind of block device) must implement two core functions: readBlock and writeBlock
- Low-level abstraction: a disk has multiple registers, including control, command block, status, and error. OS does not know such information. Only disk manufacturer knows.
- Disk manufacturer publishes a disk driver that implements readBlock and writeBlock by manipulating those registers
 - Read details in the book



Summary (Key points)

- Disk Performance (how to calculate seek time, rotational time, disk transfer time, as well as the total/average time, and throughput)
- Disk Scheduling (FCFS, SSTF, SCAN, CSCAN, SPTF)
- Polling vs. Interrupt
- Programmed I/O vs. DMA

