

Instructor: Luan Duong, Ph.D.

Electronic Submission: **11:59 pm, Wednesday, April 16, 2025**

Point: 25 points (**5%** of total grade) + Extra 5 points (1% of total grade)

Question 1: Disk Scheduling [15 points]

A) For each of the following scheduling algorithms: **SSTF**, **SCAN**, **C-SCAN**: please describe how each of them works in 1 to 2 sentences. The description should be precise enough to distinguish each algorithm from the others. [3 points]

B) Disk Scheduling [12 points]: Recall that the tracks (or so-called cylinders) of the hard drive is as described in Figure 1(a). The first track (track 0) is the outer-most track, and the inner most track is the highest track that a hard disk has. For each request, the disk head needs positioning itself onto the requested track, following some order (given by the algorithm).

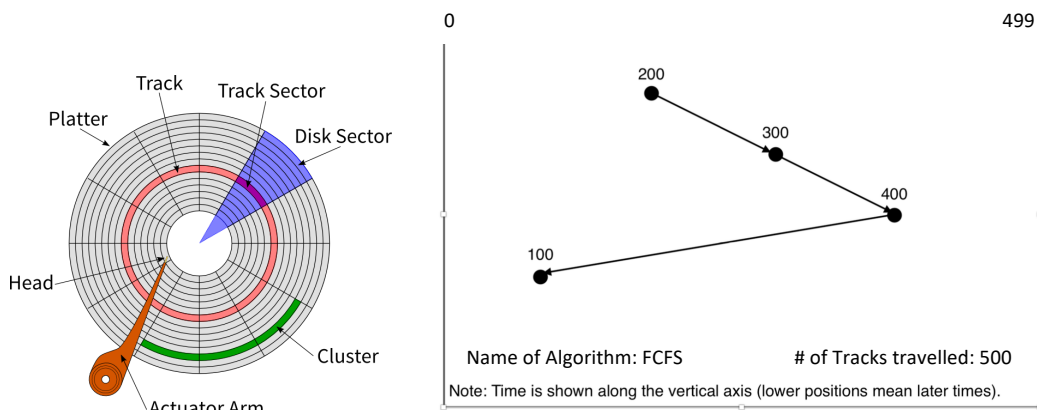


Figure 1(a): Hard Disk illustration

Figure 1(b): Example of FIFO Algorithm

Figure 1(b) shows track seeking for a hard drive with 500 tracks, numbered 0 to 499, in which the FCFS algorithm is used. The figure shows service for the following hard disk track requests: 200, 300, 400, 100. Time flows topdown. We assume that no other track requests arrive while this algorithm executes.

In this question 1B, we have a **different hard disk with 700 tracks**, numbered 0 to 699. The disk head **is at track 250 (moving toward track 699)**, and we have the following queue of hard-disk track requests: **40, 100, 150, 230, 360, 420, 500, 650**.

Please draw diagrams (as in Figure 1b) showing the movement of the head over time for SSTF and SCAN algorithm. Also, please calculate the total movement of the disk head (in # of tracks). Assume that no other track requests arrive while this algorithm executes.

1. SSTF [5 points] + Total movement of the disk head [1 point]

Name of Algorithm: _____Number of Movements: _____

2. SCAN [5 points] + Total movement of the disk head [1 point]

Name of Algorithm: _____Number of Movements: _____

Question 2: RAIDs [4 points]

Match each RAID level with its primary disadvantage. Each disadvantage should be used only once. Choose the best option for each:

- | | |
|--------------|--|
| _____ RAID-0 | a. Wastes disk capacity |
| _____ RAID-1 | b. Complicated calculation of data and parity location |
| _____ RAID-4 | c. Failure of one disk causes loss of data |
| _____ RAID-5 | d. Parity disk is performance bottleneck |

Question 3: Deadlocks [6 points]

Suppose a system has four processes, $P = \{P1, P2, P3, P4\}$ and three resource types, $R = \{R1, R2, R3\}$. In addition, assume we have two instances of R1, two instances of R2, and two instances of R3. The current state of the system is defined by the following requests and assignments:

- P1 requests an instance of R1;
- P2 requests an instance of R2;
- P3 has been assigned an instance of R1 and P3 requests an instance of R2 and an instance of R3;
- P4 has been assigned an instance of R3 and requests an instance of R1.

Please draw a resource-allocation graph to the above requests and assignments. Is there a deadlock in this situation? You need to explicitly say “Yes” or “No”. If you conclude there is no deadlock, your justification must include a sequence of execution for the processes showing that all processes can execute to completion. If you conclude there is deadlock, describe when deadlock could happen? (Note that in this situation, there is only one conclusion: deadlock or not deadlock)

Question 4: iNode [Extra 5 points] [Referred to lecture note, Lecture_Topic_8b, slide 62-68]

Consider a file system (UNIX-based) with 12 direct pointers, 1 indirect pointer, 1 double-indirect pointer, and 1 triple indirect pointer in **the i-node**. Disk blocks' size is 8K bytes, and each pointer pointing to a disk block needs 4 bytes.

- a. With this design, what is the largest supported file? You do not need to calculate the final numerical value, but you need to show the full expression. (The multiplication expression) [2 points]
- b. How many disk reads required to **read block 14** of the file named **/a**? Assume that nothing relevant is in the file-cache (e.g., no i-nodes and no data blocks) and that the root directory is only one-block long. Please describe each disk read in details. (i.e. Read [content] from where to get where) [3 points]