Lab 3 – Gage Farmer.308

<u>Experiment 1</u>
Running parallel1 always gets the same result as the sequential version, and when ran 20 times in a row it got the same result each time.

<u>Experiment 2</u>
Running parallel2 got the same result each time as parallel1 and sequential. When ran 20 times in a row the answer came out unchanged.

<u>Experiment 3</u>
An improvement over parallel2 I implemented in parallel3 was the removal of mutex locks during the calculation of column sums. This was because the threads shouldn't use the same data as eachother and should cut the lock use in half.